



MINERVA

PARALLEL AI



WHY PARALLEL TOOLS IN AI?





MEMORY OCCUPANCY

MODEL STATES

● Parameters

● Gradients

● Optimizer states (such as momentum and variances in Adam)

[1] <https://arxiv.org/pdf/2101.06840>

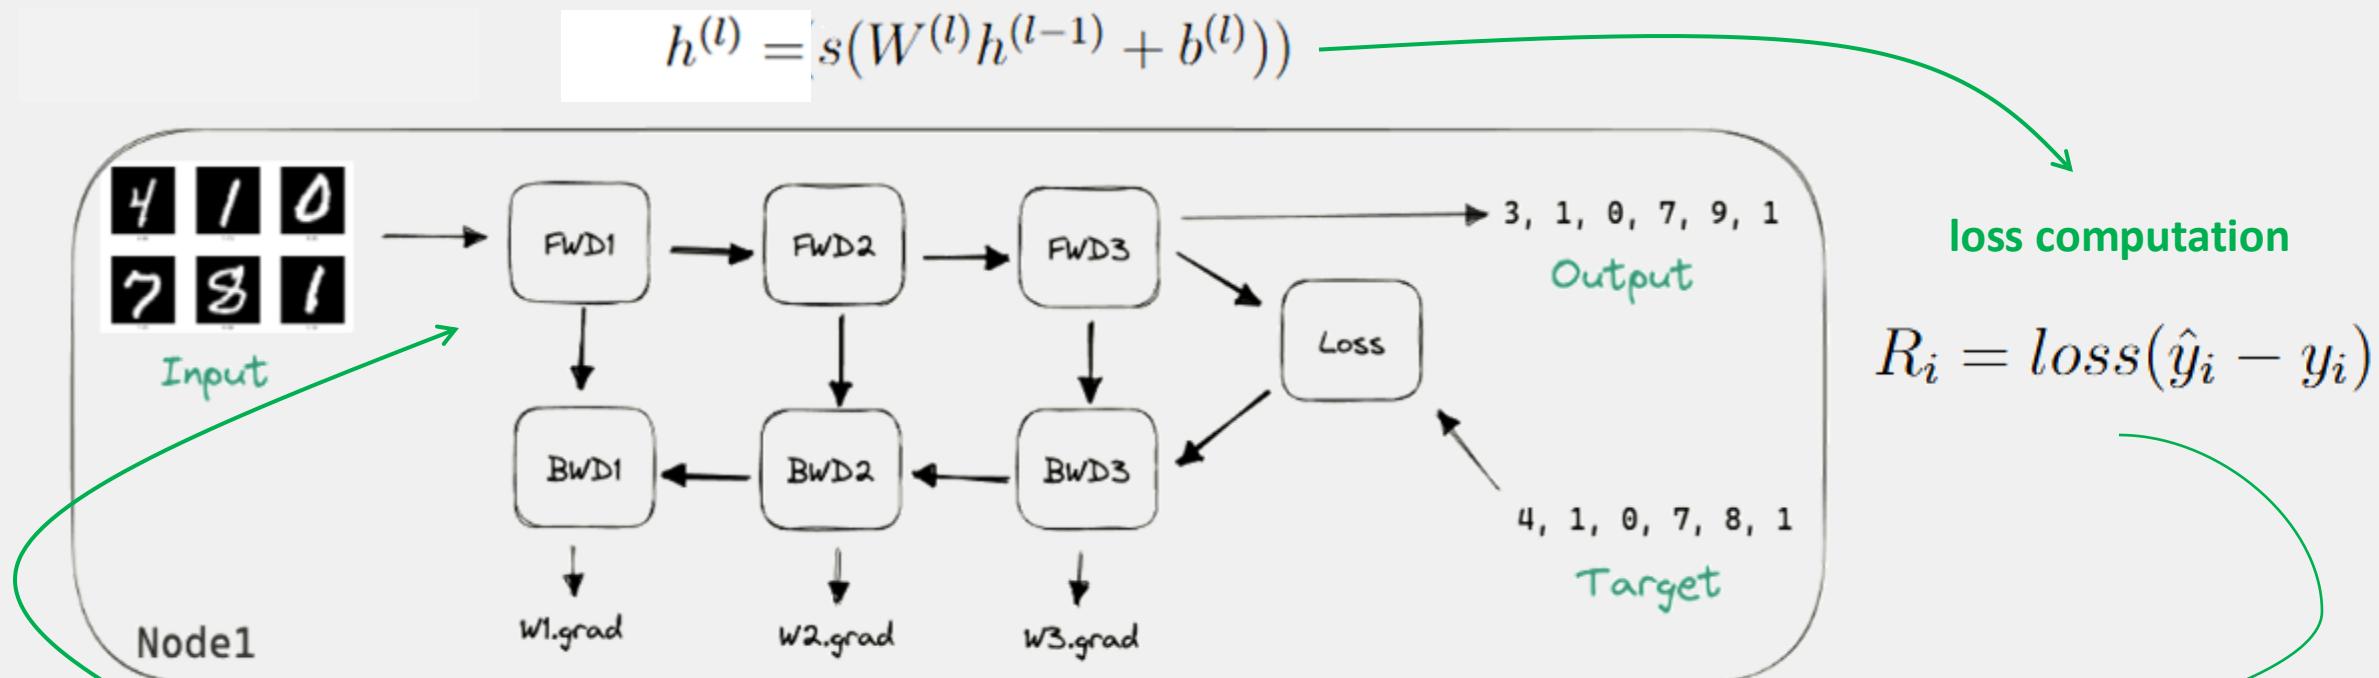
RESIDUAL STATES

● Activations

● Temporary buffers

● Unusable fragmented memory

BASICS OF DNN TRAINING



optimizer step

$$W^{(l)}(j, k) = W^{(l)}(j, k) - \gamma \sum_{i=k}^{k+t} \frac{\partial R_i}{\partial W^{(l)}(j, k)}$$

$$b^{(l)}(j) = b^{(l)}(j) - \gamma \sum_{i=k}^{k+t} \frac{\partial R_i}{\partial b^{(l)}(j)}$$

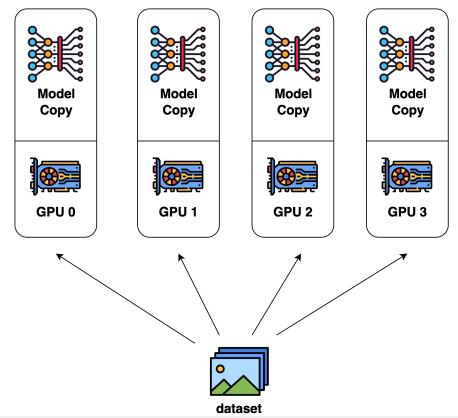
gradient computation

$$\frac{\partial R_i}{\partial b^{(l)}(j)} = \delta^{(l)}(j)$$

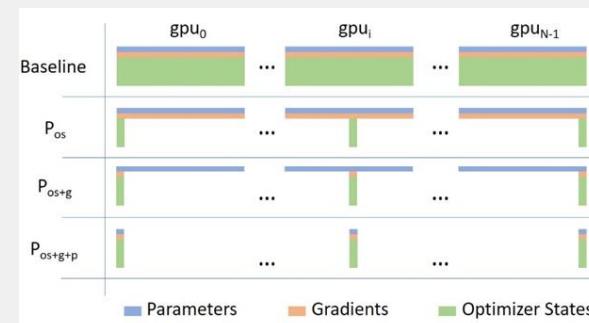
$$\frac{\partial R_i}{\partial W^{(l)}(j, k)} = \delta^{(l)}(j)h_i^{(l)}$$

DIFFERENT TYPES OF PARALLELISM

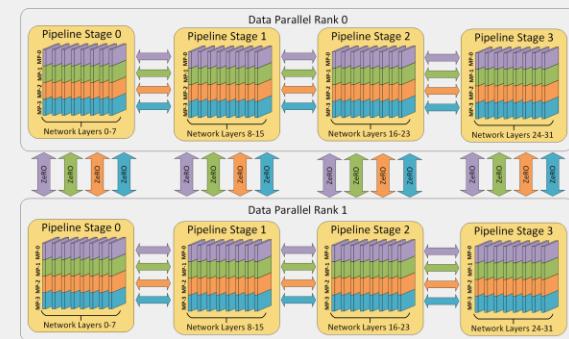
Data Parallelism



Model Sharding



Model Parallelism



https://colossalai.org/docs/concepts/paradigms_of_parallelism/

<https://www.deepspeed.ai/tutorials/pipeline/>

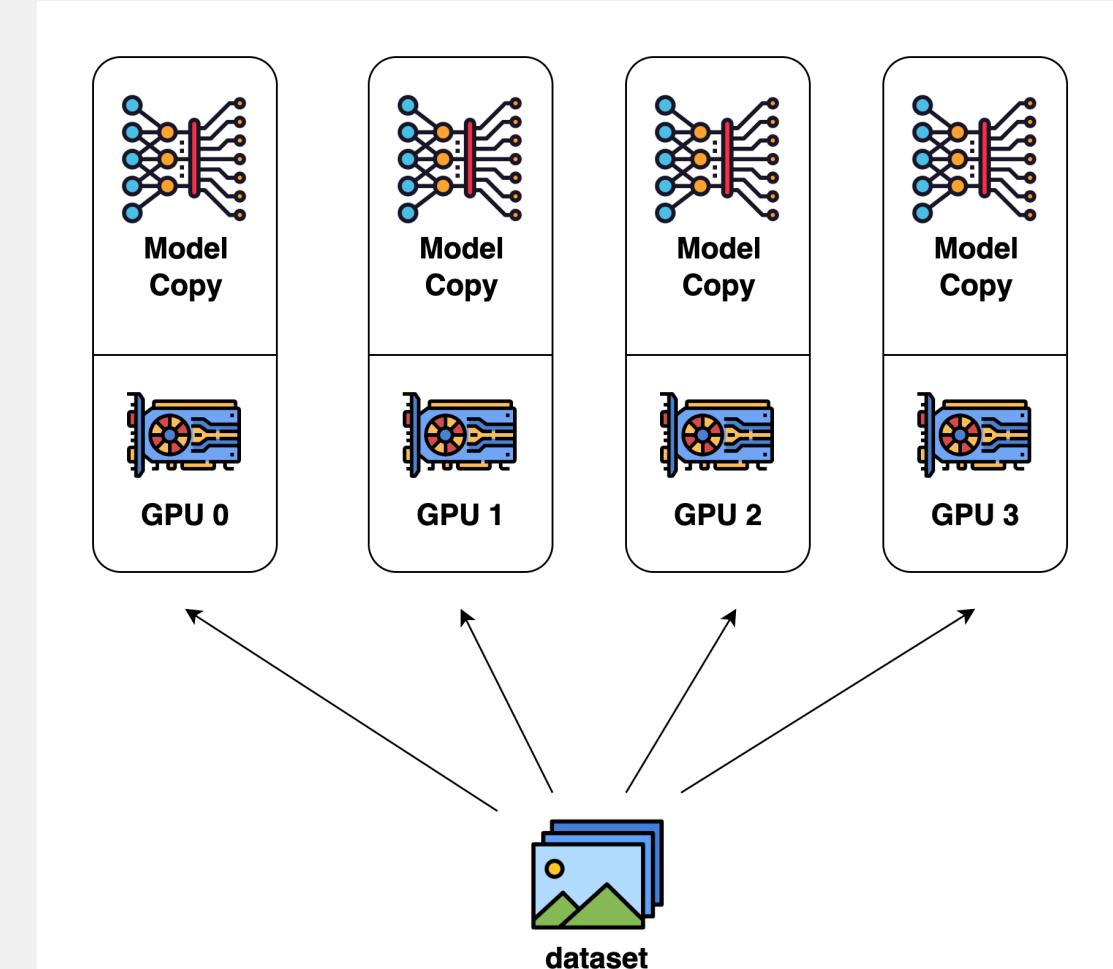


DATA PARALLELISM



HOW IT WORKS

- Each device holds a **full copy** of the model

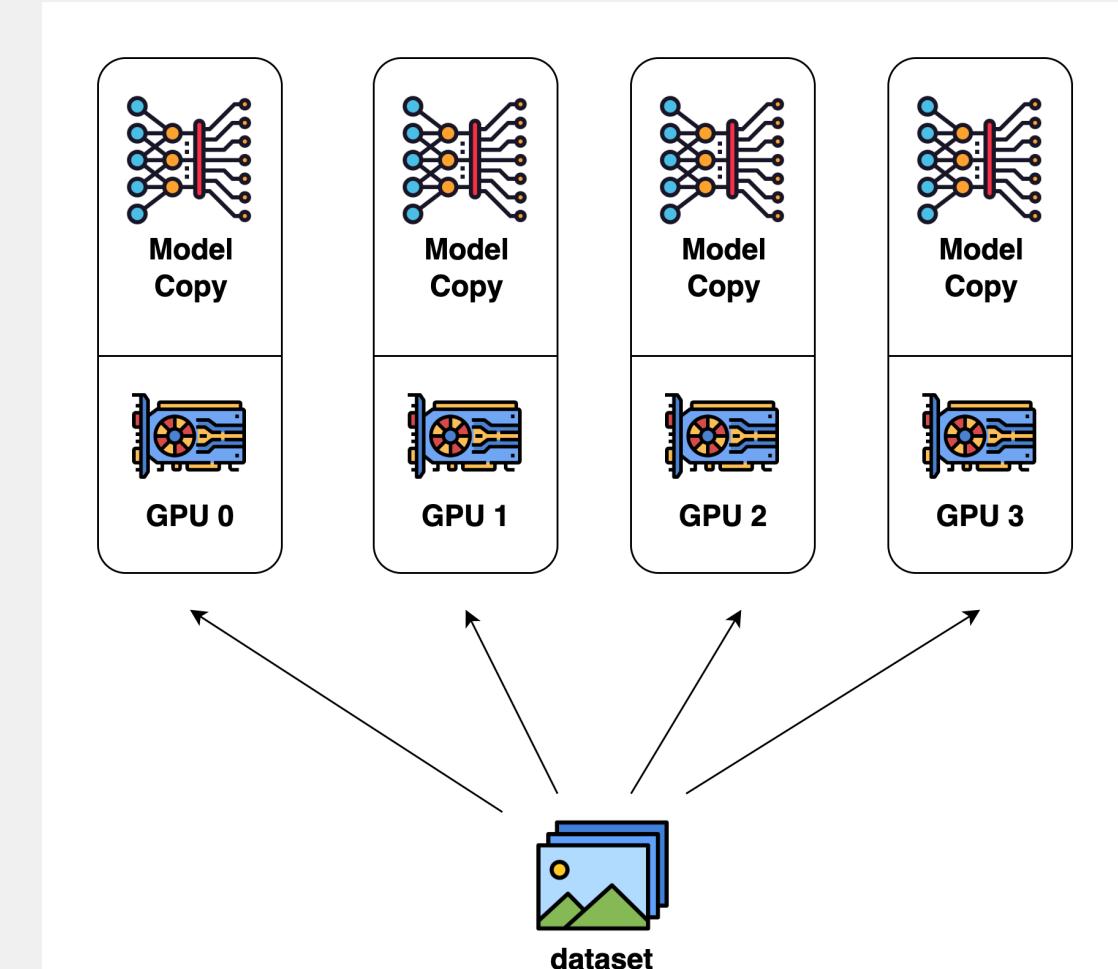


https://colossalai.org/docs/concepts/paradigms_of_parallelism/



HOW IT WORKS

- Each device holds a **full copy** of the model
- Data are split on multiple GPUs (single or multi-node)



https://colossalai.org/docs/concepts/paradigms_of_parallelism/

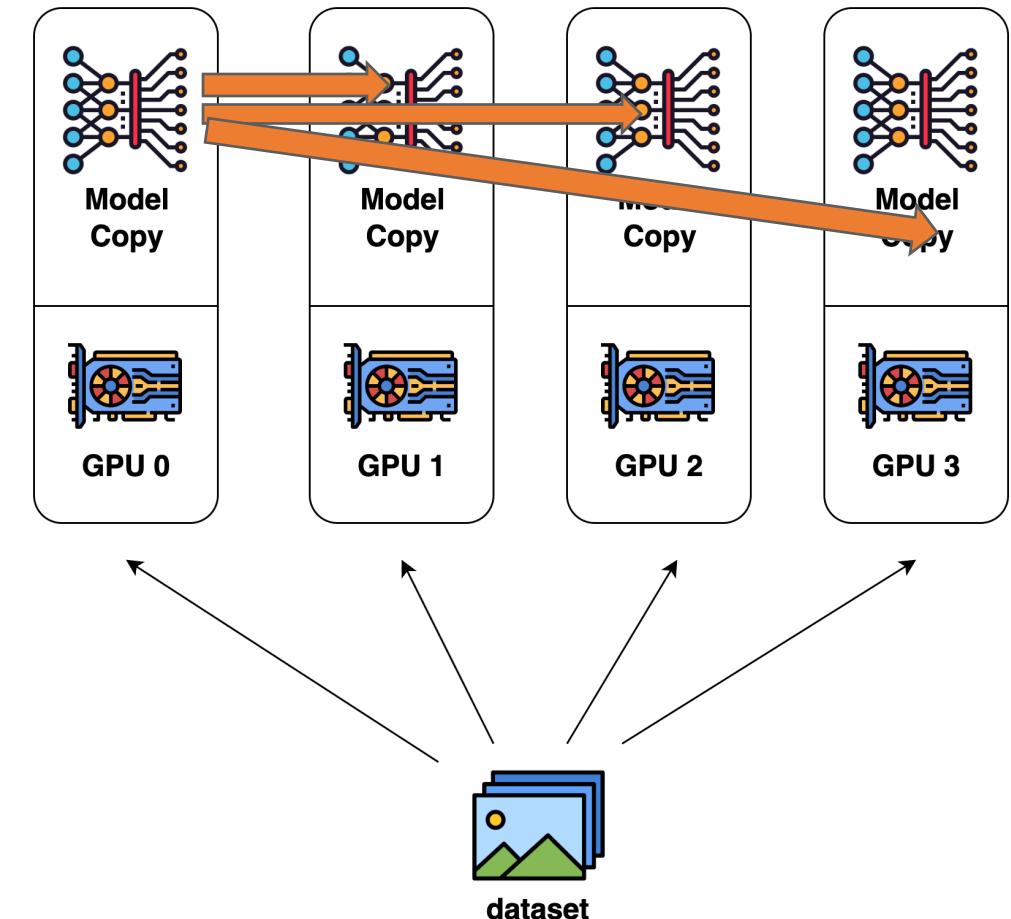


HOW IT WORKS

- Each device holds a **full copy** of the model
- Data are split on multiple GPUs (single or multi-node)
- After some backpropagation steps, **gradients** are **synchronized** to avoid convergence issues

https://colossalai.org/docs/concepts/paradigms_of_parallelism/

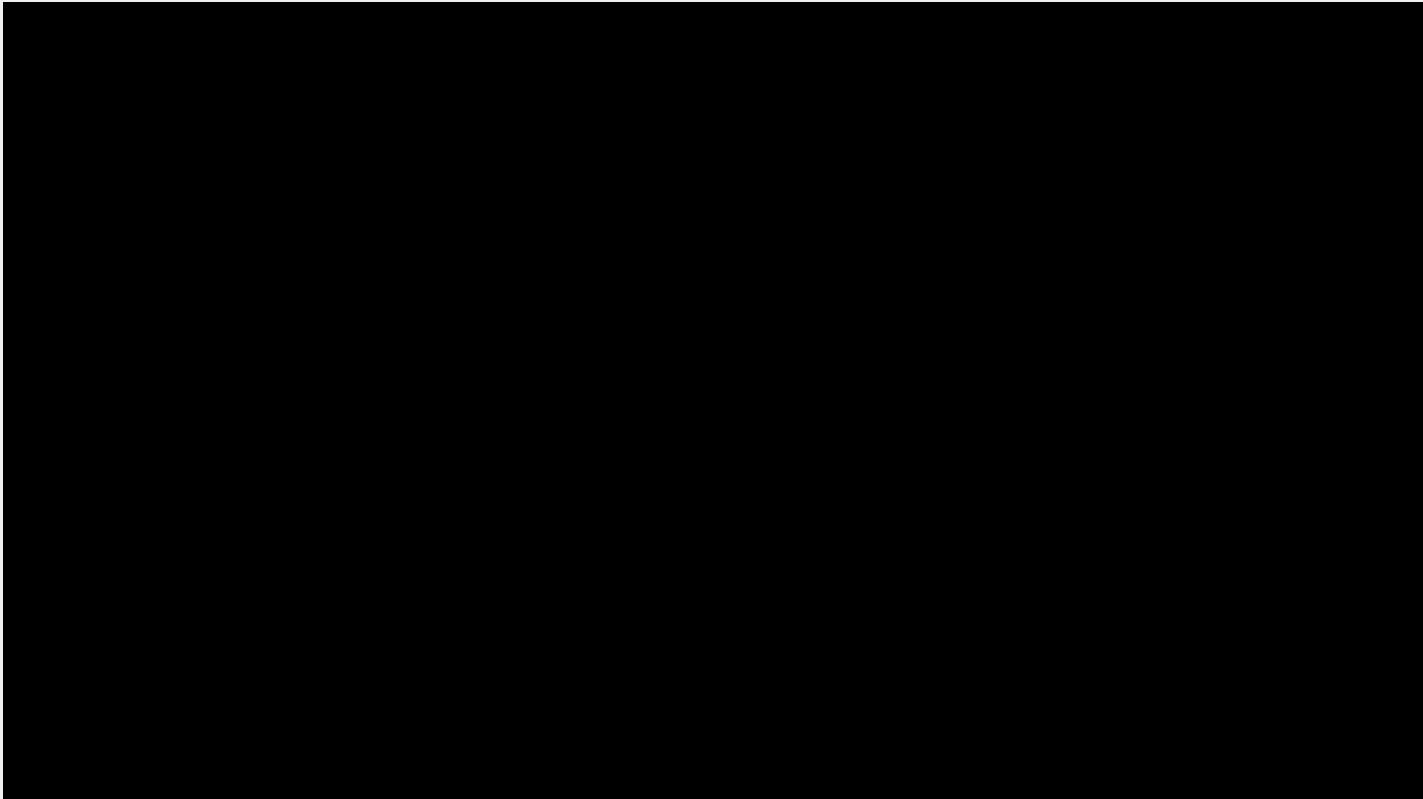
Ring ALL-REDUCE





HOW IT WORKS

Data parallelism enables distributed training by **communicating gradients before the optimizer step** to make sure that parameters of all model replicas are updated using exactly the same set of gradients, and hence model replicas can stay consistent across iterations.



<https://docs.nvidia.com/nemo-framework/user-guide/24.07/nemotoolkit/features/parallelisms.html>

● **DataParallel** works on a **single** multi-GPU **node**, slower than DistributedDataParallel due to GIL contention

● **DistributedDataParallel** works on a **multi-node** multi-GPU setting, it can be combined also with **model parallelism**



[1] https://pytorch.org/tutorials/intermediate/ddp_tutorial.html

[2] <https://towardsdatascience.com/distributed-parallel-training-data-parallelism-and-model-parallelism-ec2d234e3214>

[3] <https://medium.com/codex/a-comprehensive-tutorial-to-pytorch-distributeddataparallel-1f4b42bb1b51>

DISTRIBUTED DATAPARALLEL (DDP)

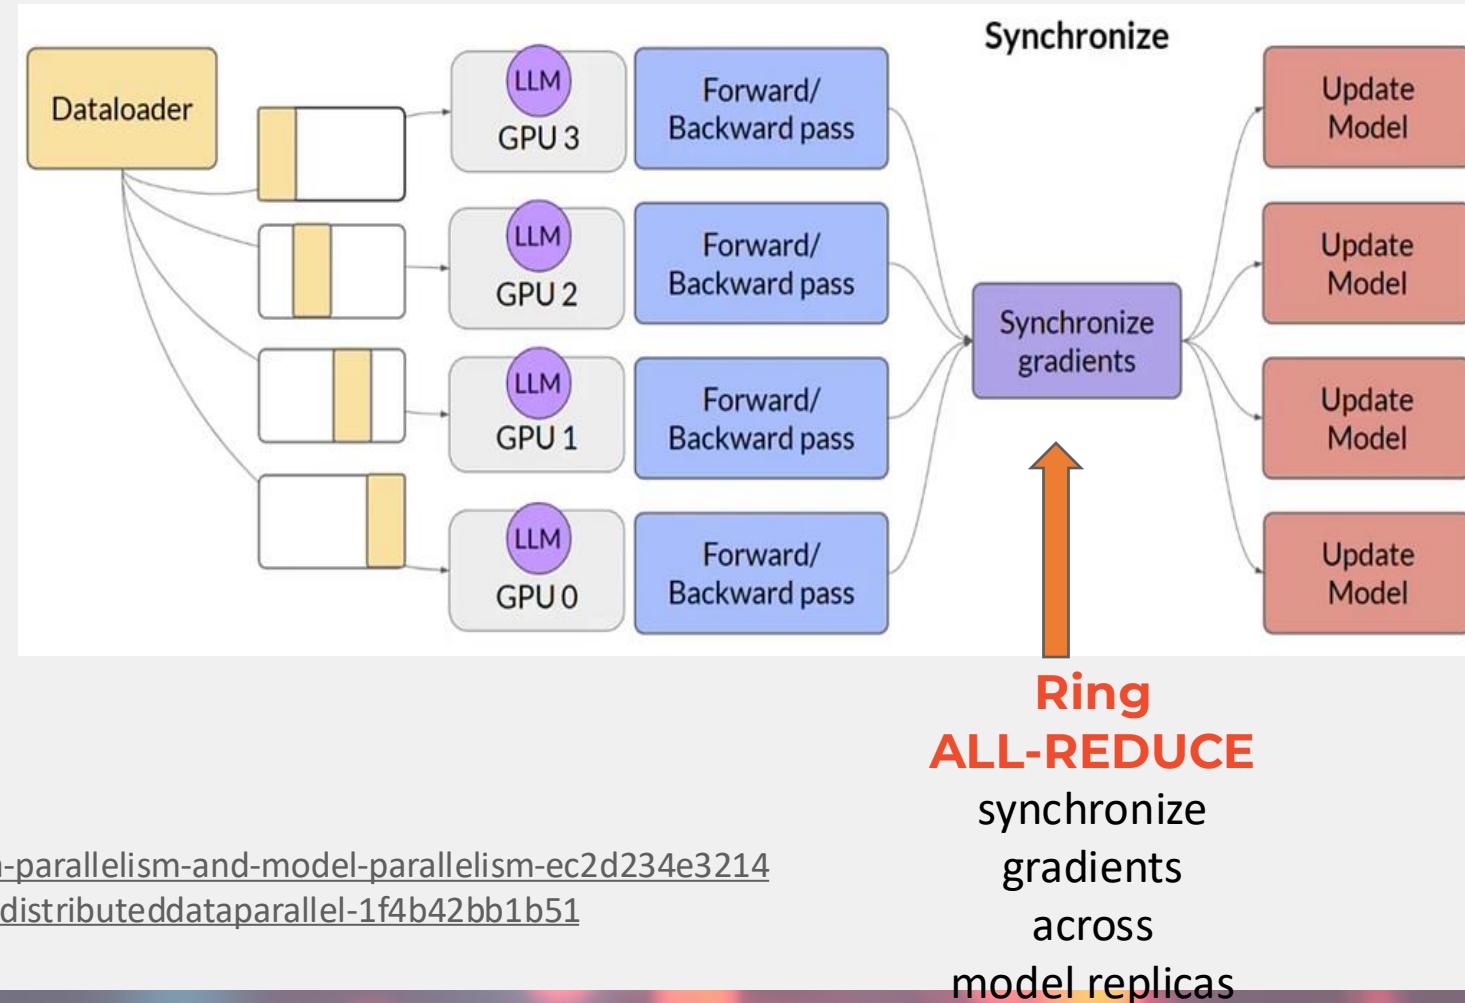
Faster and more flexible than DP, but does **not split the data** across multiple GPUs

Setup the process group
(always the same)

Split the data across multiple GPUs
(`torch.utils.data.DistributedSampler`)

Wrap the model with DDP

Nccl backend is highly recommended
(both for single-node and multi-node)



[1] https://pytorch.org/tutorials/intermediate/ddp_tutorial.html

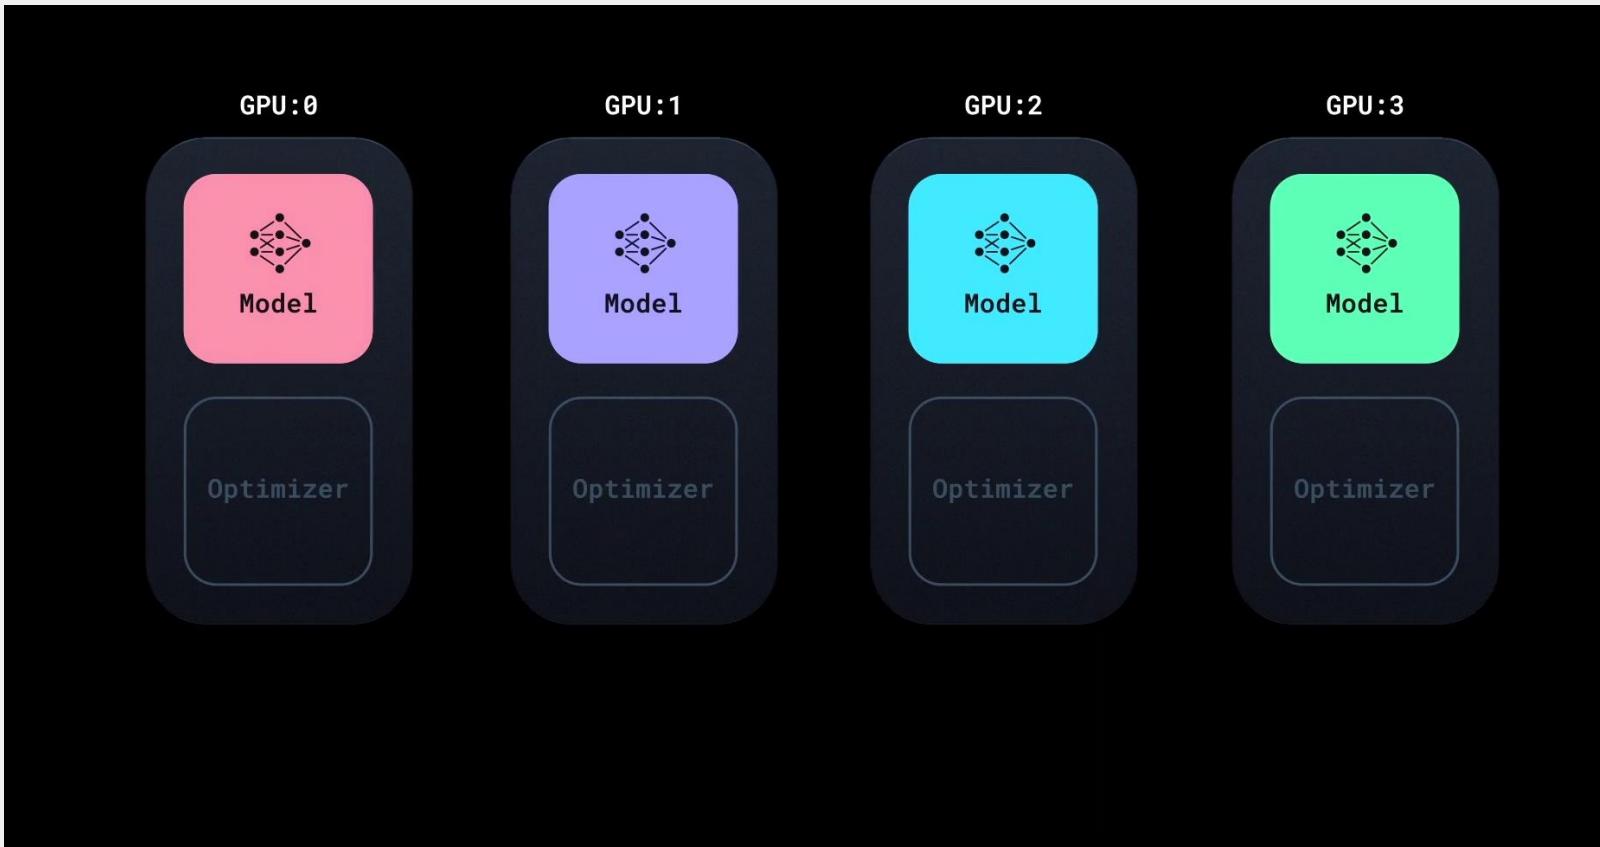
[2] <https://towardsdatascience.com/distributed-parallel-training-data-parallelism-and-model-parallelism-ec2d234e3214>

[3] <https://medium.com/codex/a-comprehensive-tutorial-to-pytorch-distributeddataparallel-1f4b42bb1b51>

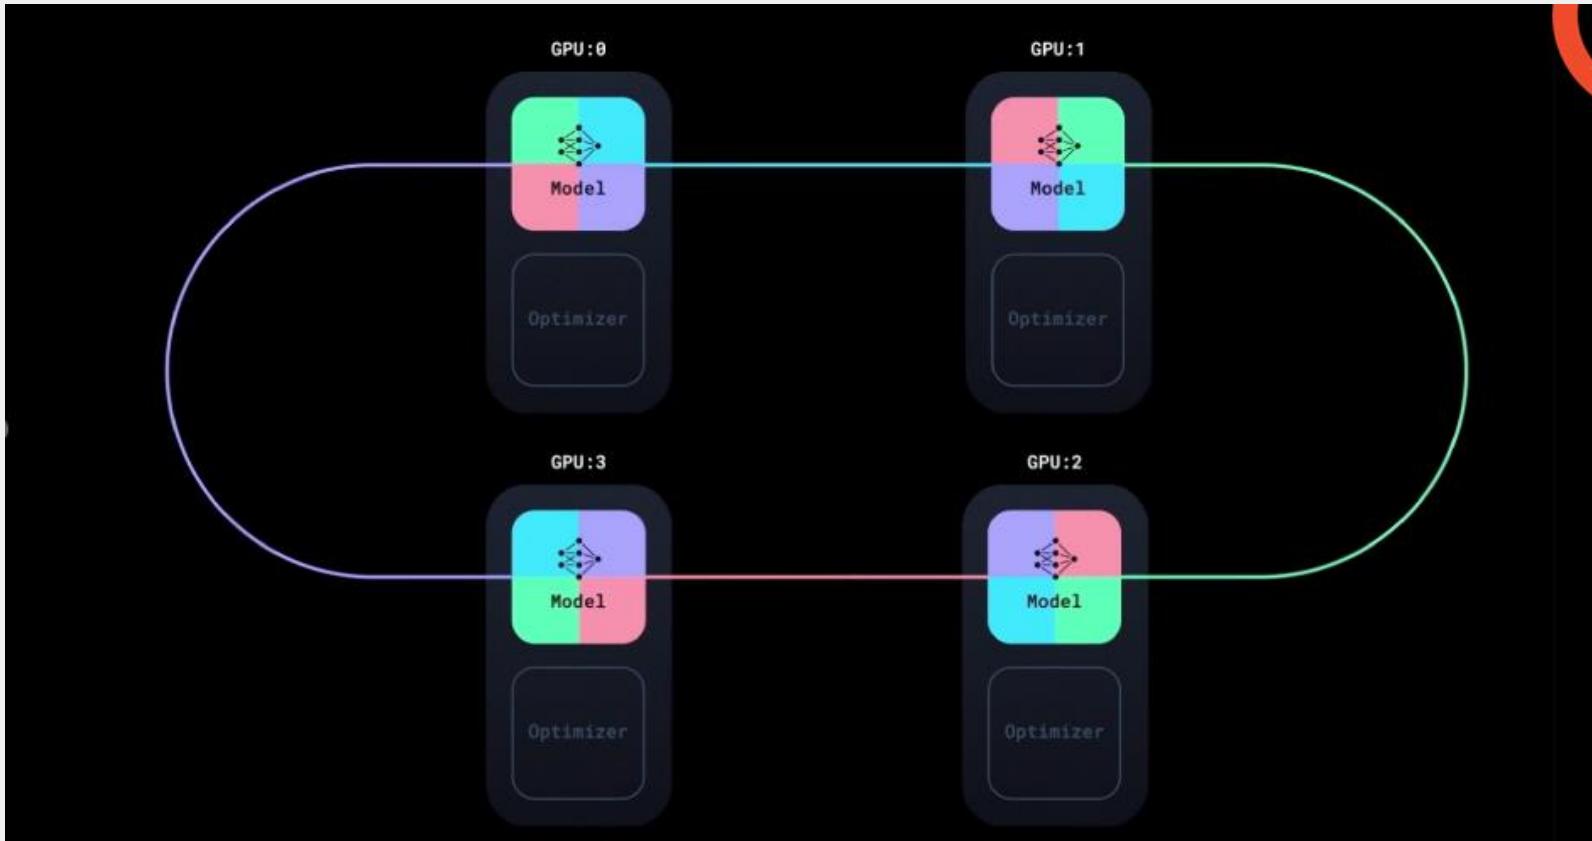
[4] <https://arxiv.org/pdf/2006.15704>



DISTRIBUTED DATAPARALLEL (DDP)



DISTRIBUTED DATAPARALLEL (DDP)



**Ring
ALL-REDUCE**
synchronize
gradients
across
model replicas



ACCELERATE – HOW TO LAUNCH LLM TRAINING

To implement DDP using HF Accelerate, you just need to follow these 3 simple steps:

1. Initialize Accelerate processes in our training
2. Define a .yaml config file
3. Launch it with accelerate launch





ACCELERATE

Distributed training across multiple nodes with a simple configuration files and few changes in the python script

Several level of parallelism: **data parallelism**, model sharding, hybrid sharding. Handles both multi CPUs and multi GPUs using Only a config file!

```
compute_environment: LOCAL_MACHINE
debug: true
distributed_type: MULTI_GPU
downcast_bf16: 'no'
enable_cpu_affinity: false
machine_rank: 0
main_training_function: main
mixed_precision: bf16
num_machines: 4
num_processes: 16
rdzv_backend: c10d
same_network: true
tpu_env: []
tpu_use_cluster: false
tpu_use_sudo: false
use_cpu: false
```



ACCELERATE

To launch the training with the new Accelerate configuration performing distributed training it will be enough to run

```
accelerate launch --config_file config_accelerate.yaml my_script.py <py_args>
```

In SLURM it becomes

```
accelerate launch \
--main_process_ip "$MASTER_ADDR" \
--main_process_port $MASTER_PORT \
--machine_rank $SLURM_PROCID \
--rdzv_backend c10d \
--config_file config_accelerate.yaml
my_script.py <py_args>
```



READY TO SCALE YOUR LLM TRAINING





DISTRIBUTED DATAPARALLEL (DDP)

- Copy the parallelAI folder in your \$SCRATCH
- Go in the DDP subfolder
- Complete "FFT_DDP.py" exercise using :
 - "/leonardo_work/tra26_minwincs/DATA//Bitext-customer-support-llm-chatbot-training-dataset" dataset
 - "/leonardo_work/tra26_minwincs/models/Llama-3.2-1B-Instruct" model
- Run the code with 1 node ("job_FFT_DDP.sh")
- Run the code also with 2 nodes and check the training times
- Set **#SBATCH --reservation=s_tra_minwincs** for accessing the reservation



MINERVA

MODEL SHARDING

WHY MODEL SHARDING?



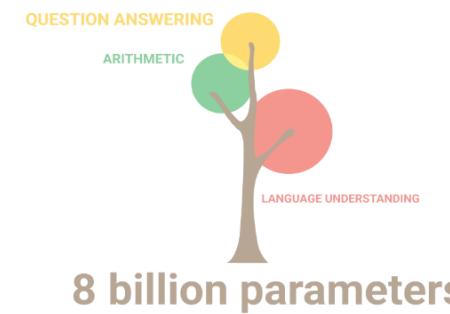
WHY MODEL SHARDING?

Full-Precision : fp32 (4 bytes)

Params	7B x 4	28 GB
Activation	7B x 4	28 GB
Gradients	7B x 4	28 GB
Optimizer	7B x 4 x 2	56 GB
TOT	140 GB	

Half-Precision : fp16 or bf16 (2 bytes)

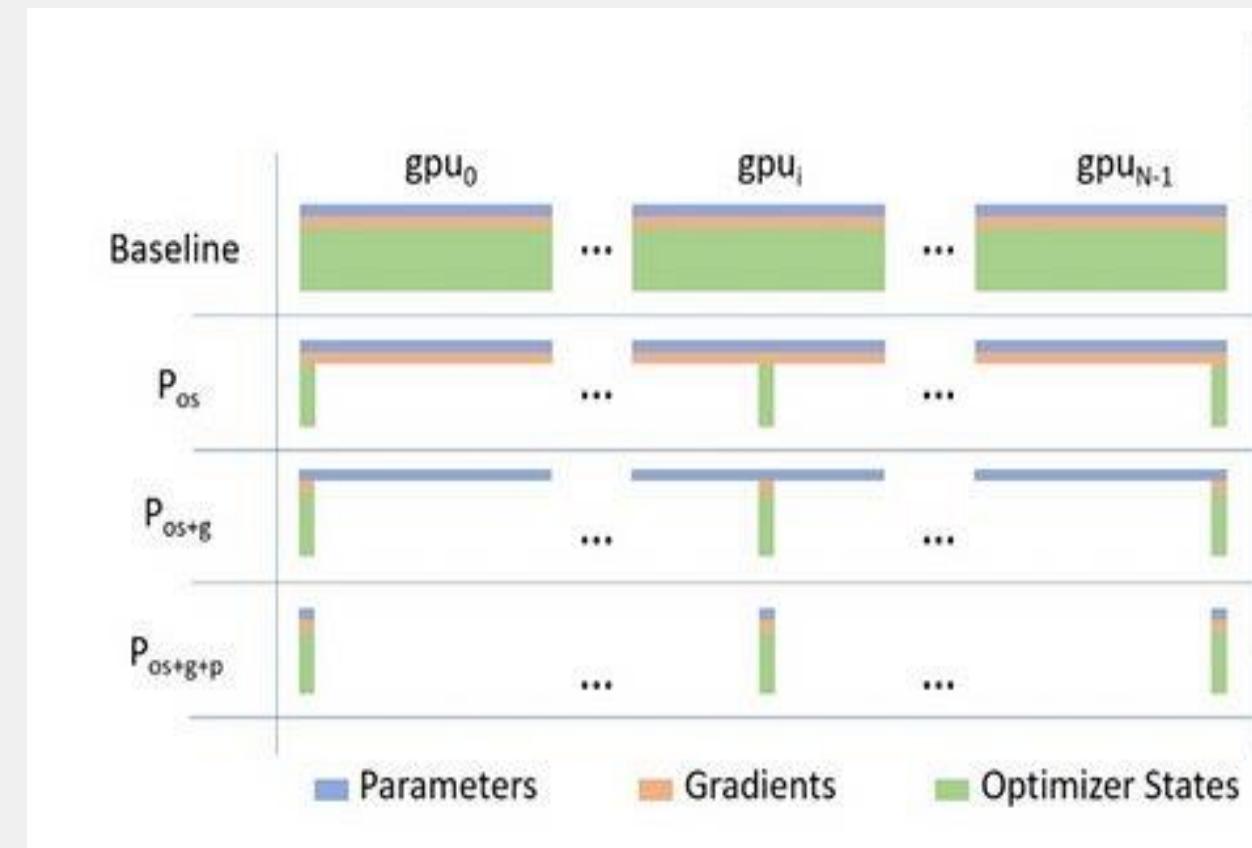
Params	7B x 2	14 GB
Activation	7B x 2	14 GB
Gradients	7B x 2	14 GB
Optimizer	7B x 2 x 2	28 GB
TOT	70 GB	



In Leonardo we have A100 GPUs with 64 GB each! → OoM error!

SHARDING: HOW IT WORKS

- Each device holds a "**shard**" of the model **parameters**, **optimizer states** and **gradients**
- Data are **split** on multiple GPUs (single or multi-node)
- Enables **larger** models **training** (the model Does not need to fit in a single gpu)
- More communications** needed wrt DDP
- Replace **All Reduce** communication With **Reduce Scatter+ All Gather**

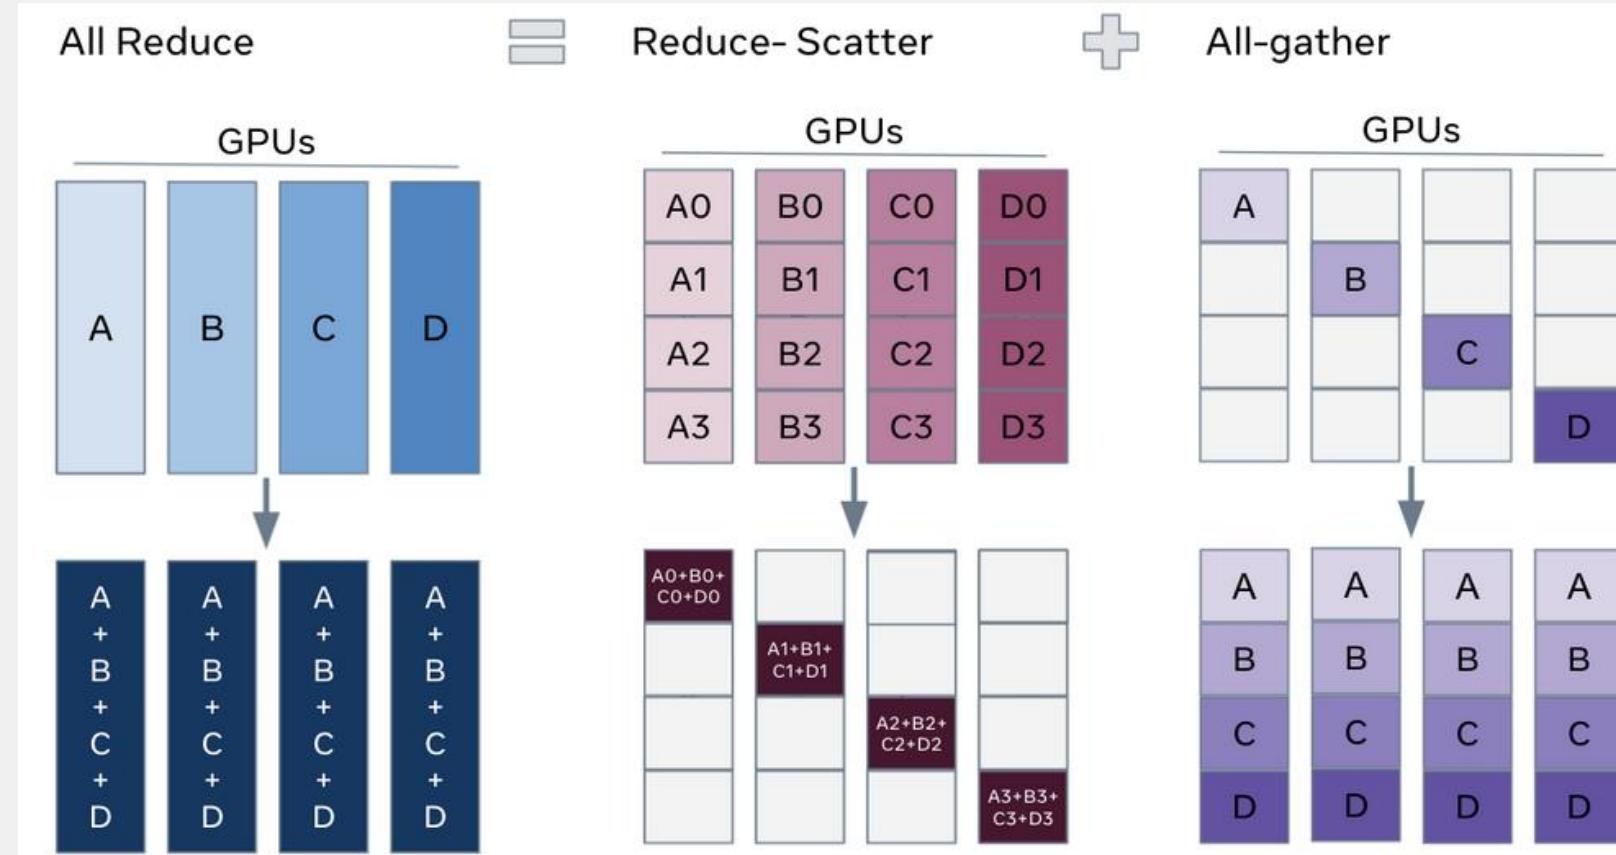


[1] <https://medium.com/@pranay.janupalli/understanding-model-sharding-and-model-parallelism-scaling-large-language-models-dee6144d0591#:~:text=%E2%80%94%20Model%20Sharding%20Usually%20involves%20a,types%20of%20operations%20across%20devices>

Reduce-Scatter:
one rank receive
the result T

All-Gather:
each rank receives
the contributions of other ranks
(but not the final result T)

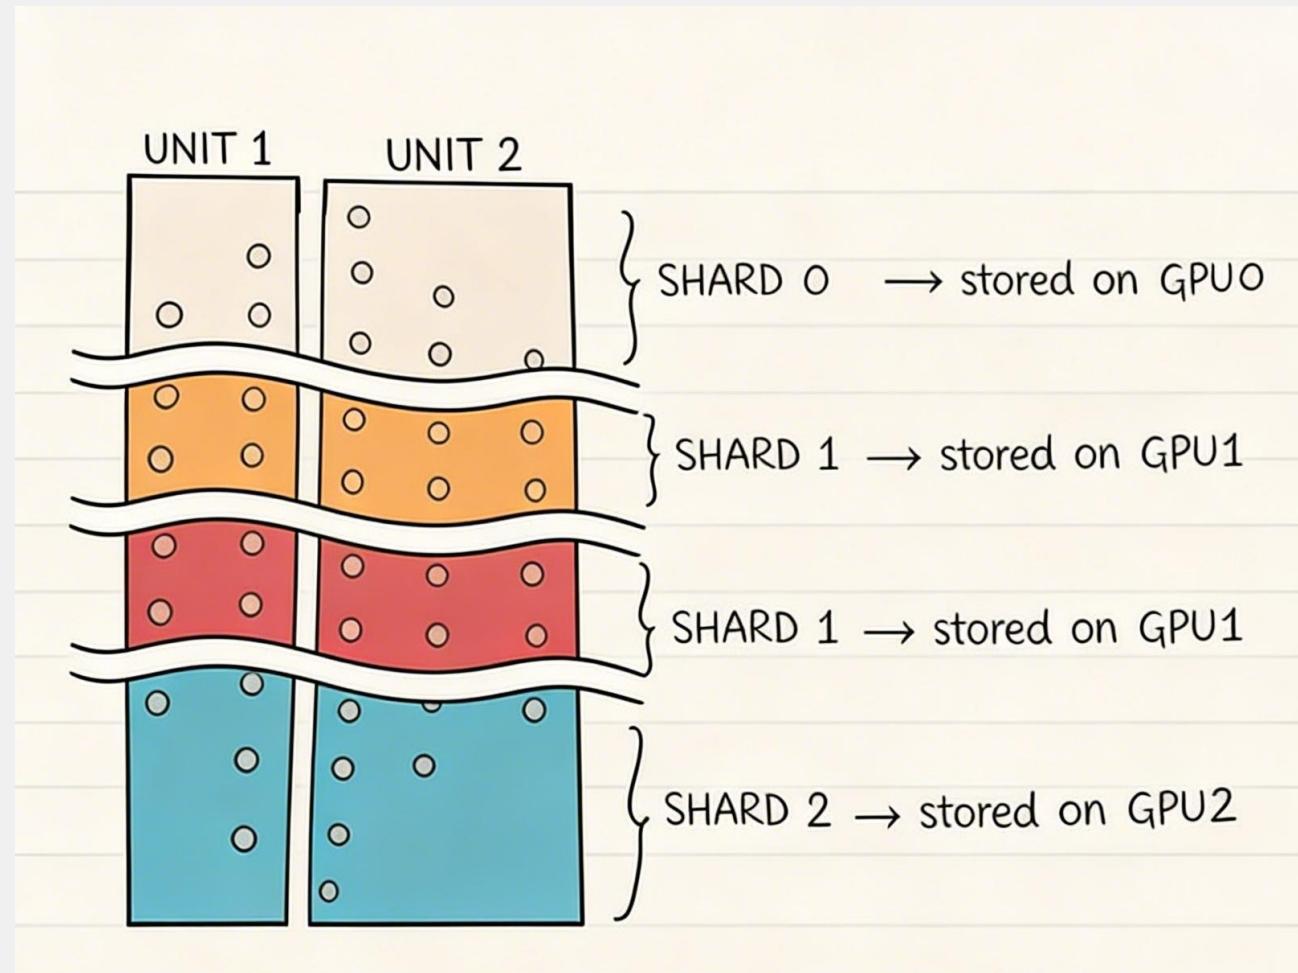
All-Reduce:
all ranks receive the result T

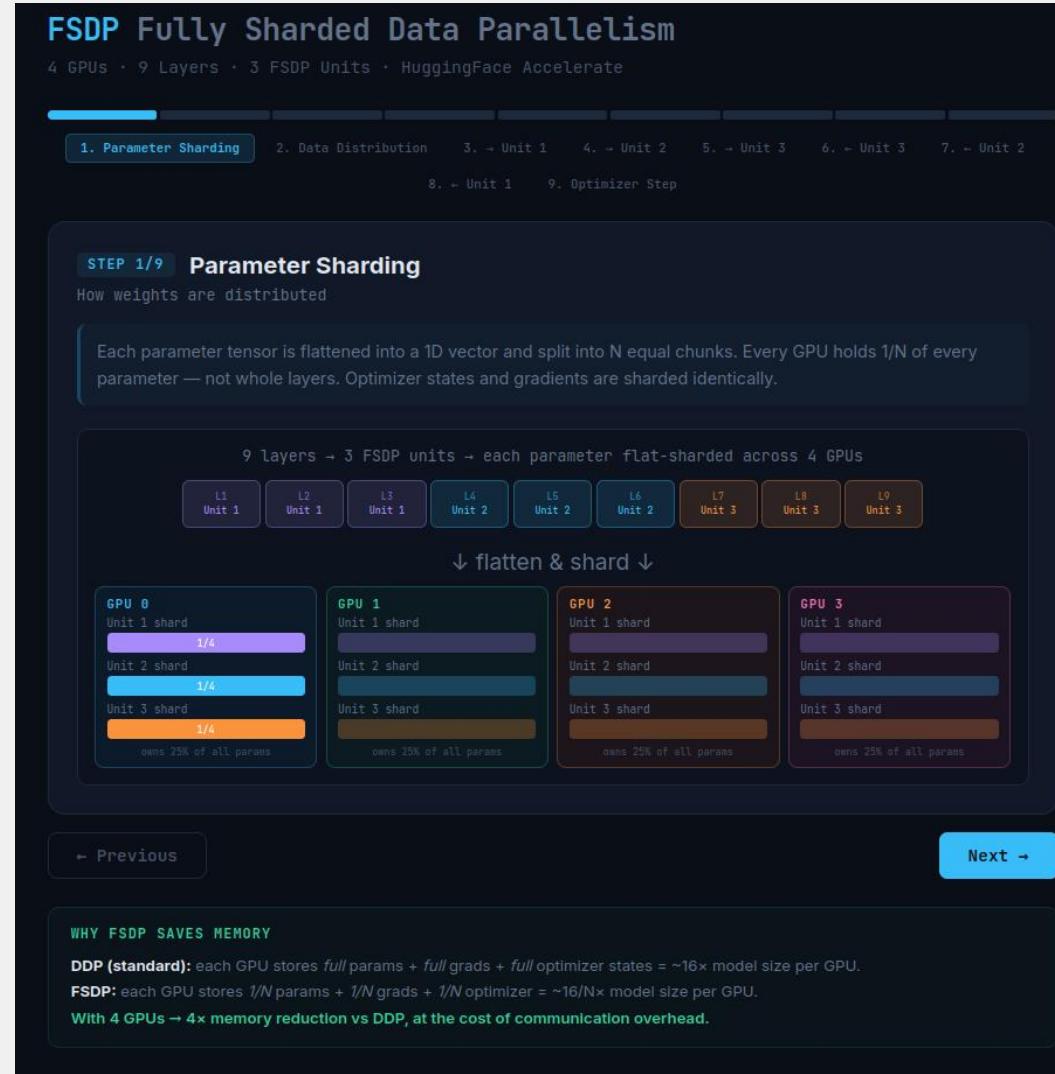


[1] https://docs.pytorch.org/tutorials/intermediate/FSDP_tutorial.html



- 4 GPUS
- 5 layers
- 2 units -> 2 all gather in the fwd







FSDP pseudocode

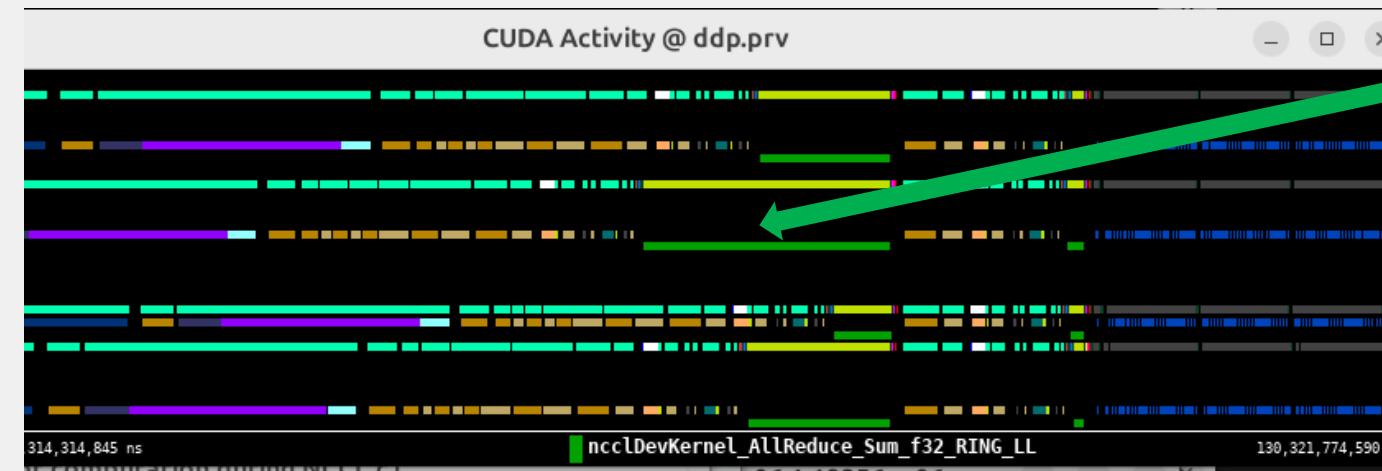
```
FSDP forward pass:  
    for layer_i in layers:  
        all-gather full weights for layer_i  
        forward pass for layer_i  
        discard full weights for layer_i  
  
FSDP backward pass:  
    for layer_i in layers:  
        all-gather full weights for layer_i  
        backward pass for layer_i  
        discard full weights for layer_i  
        reduce-scatter gradients for layer_i
```

<https://engineering.fb.com/2021/07/15/open-source/fsdp/>



DDP VS FSDP

CUDA Activity @ ddp.prv



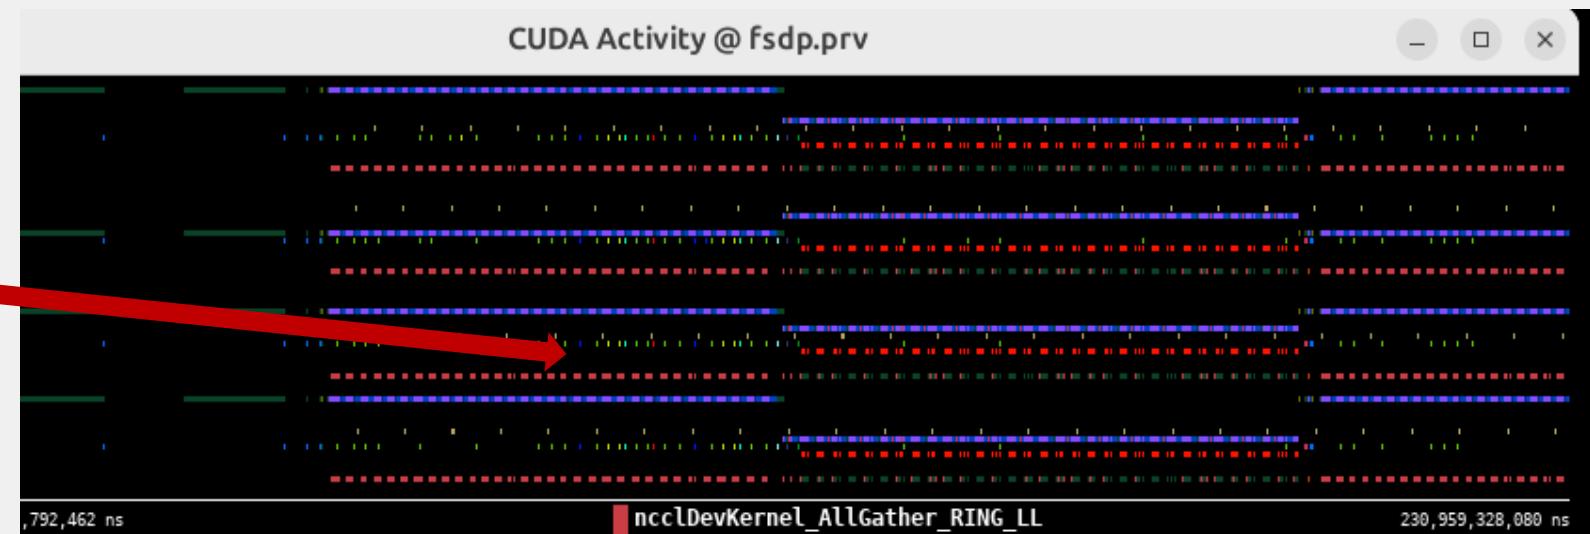
DDP
Bwd: All_reduce

FSDP

Fwd: All_Gather

Bwd: All_Gather + Reduce Scatter

CUDA Activity @ fsdp.prv





Hybrid shared data parallelism (HSDP)

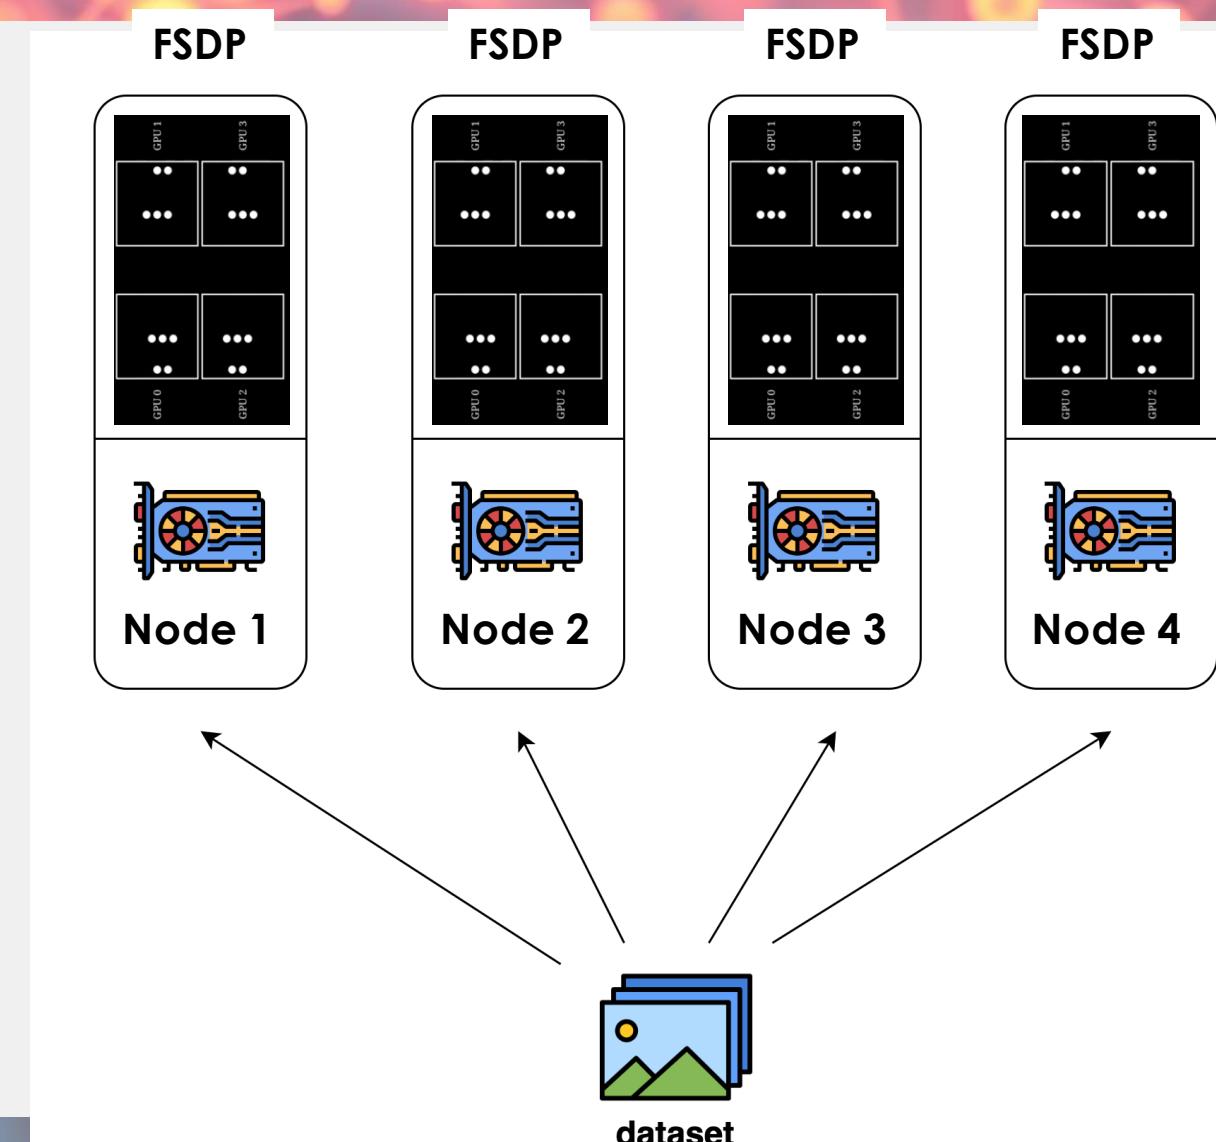
- **HYBRID SHARD:** Apply FULL SHARD **within one node**, and **replicate** parameters across nodes.

Pro:

- **Reduced** communication volume

Cons:

- **More** GPU memory needed w.r.t. FSDP





TOOLS FOR MODEL SHARDING

- 1 Accelerate
- 2 DeepSpeed
- 3 Olmo-core
- 4 Modalities



FSDP EXERCISE

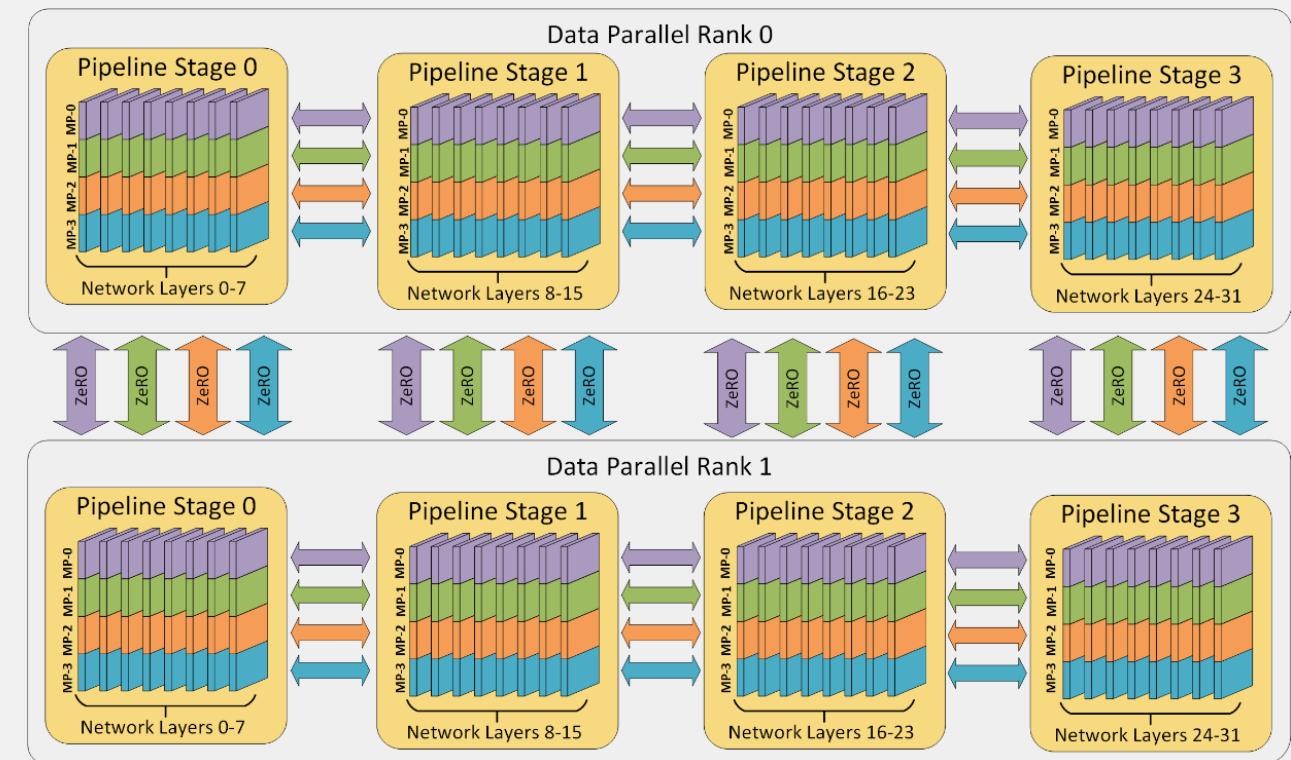
- Go in the FSDP subfolder
- Complete "FFT.py" exercise
- Complete "config_FFT_FSDP.yaml"
- Complete the jobsript "job_FFT.sh"
- Run the code with 2, 4,8 GPU
- Set **#SBATCH --reservation=s_tra_minwinc** for accessing the reservation



MODEL PARALLELISM

MODEL PARALLELISM

Model Parallelism is a technique used to split a neural network across multiple GPUs or nodes where **different GPUs handle different parts of the model.**

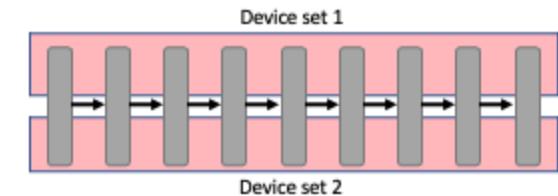


MODEL PARALLELISM

HORIZONTAL SPLIT – TENSOR PARALLELISM

Each tensor is split up into multiple chunks saved on different gpus. During processing each shard gets processed separately. Then, the results are synchronized at the end of the step.

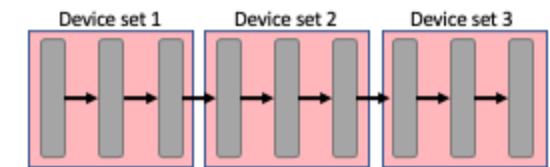
Tensor Parallelism (TP)



VERTICAL SPLIT - PIPELINE PARALLELISM

The model is split at layer level on multiple gpus (i.e. only one or several layers are stored in the same GPU). To avoid big pipeline bubbles each gpu process a portion of the batch (micro batch)

Pipeline Parallelism (PP)

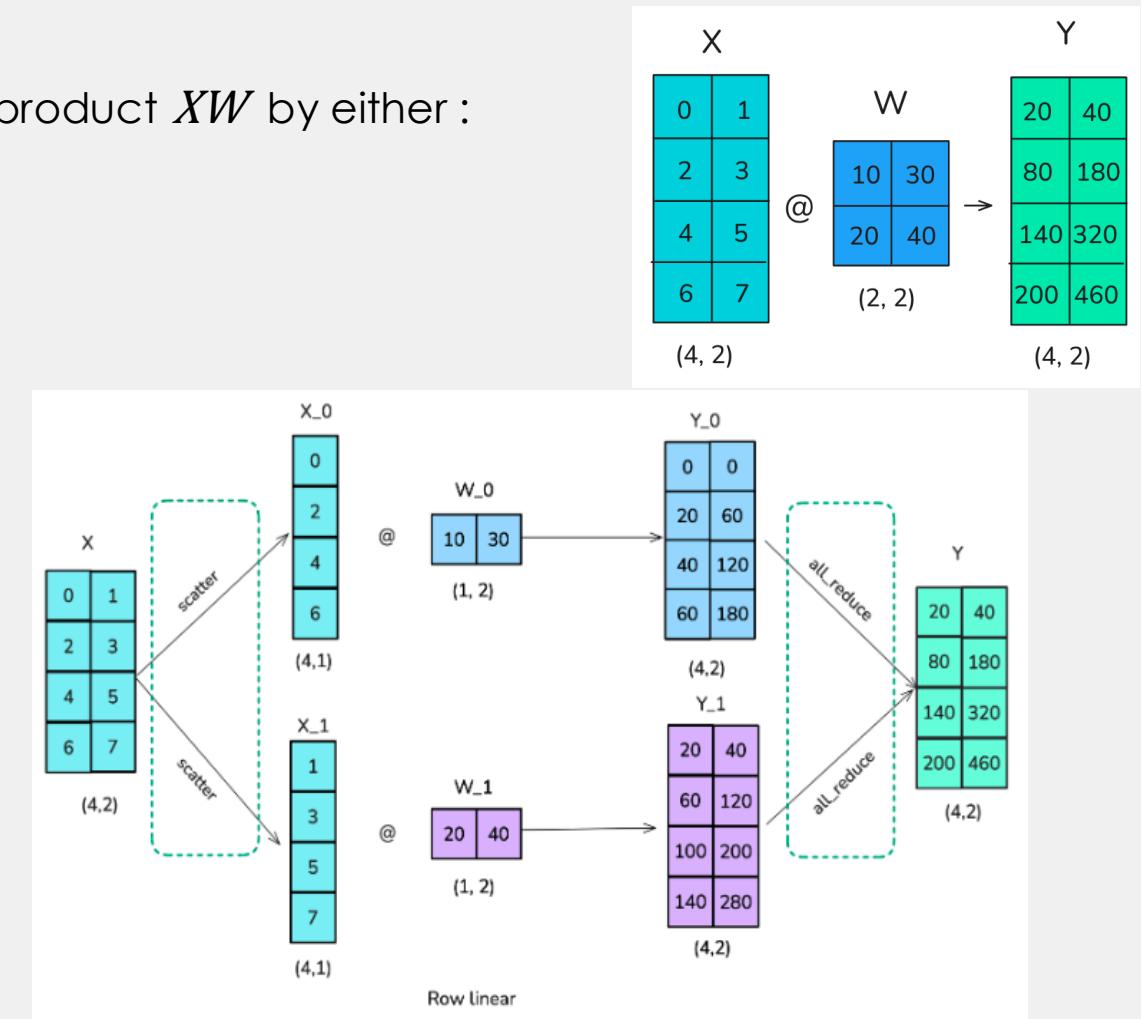
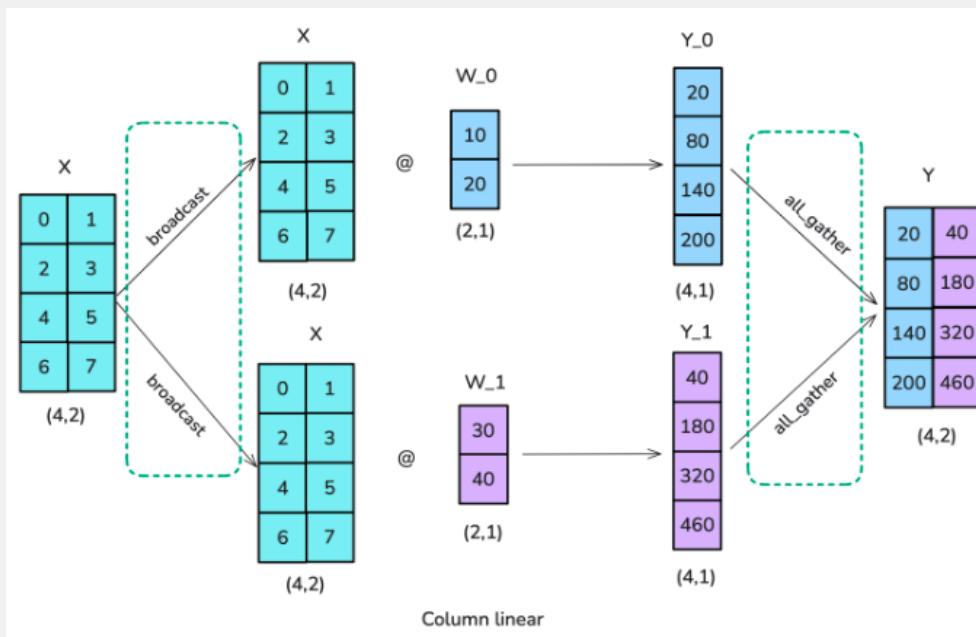


TENSOR PARALLELISM – HOW IT WORKS

Given two matrices X and W , we can compute the matrix product XW by either :

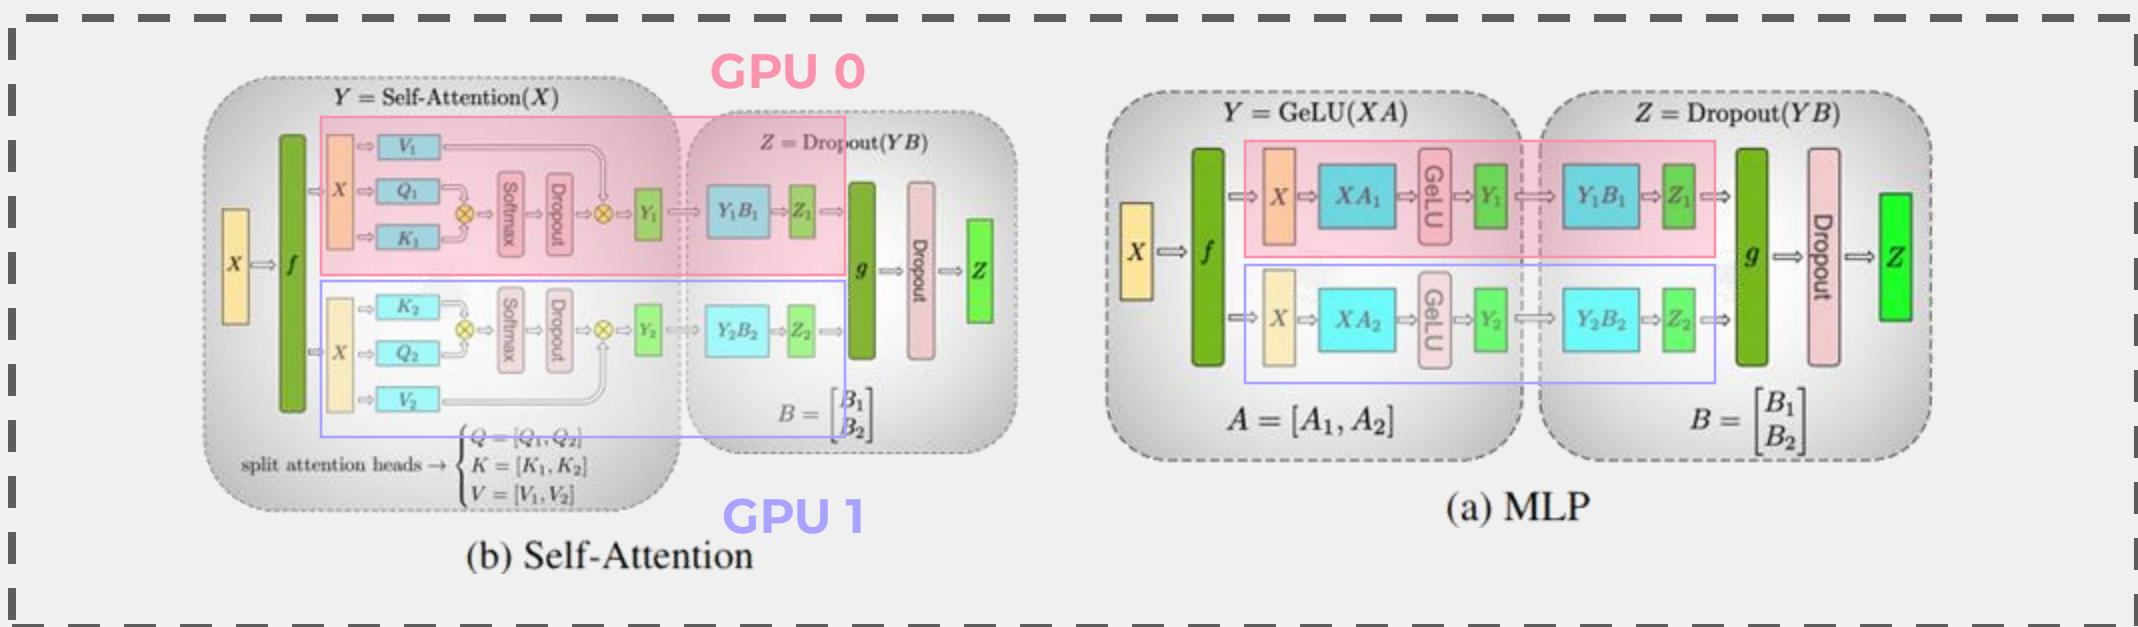
- multiplying each **column** of W individually
- multiplying each **row** of W individually

and combining the results



TENSOR PARALLELISM – HOW IT WORKS

A Transformer model is made of two main building blocks:
 a feedforward multi-layer perceptron (MLP) block and a multi-head attention block.
 We can apply tensor parallelism to both.

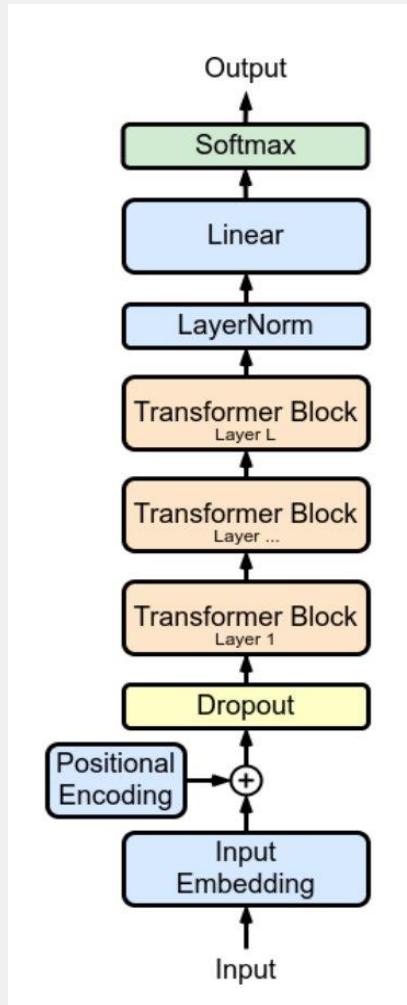


[1] <https://github.com/huggingface/transformers/issues/10321#issuecomment-783543530>

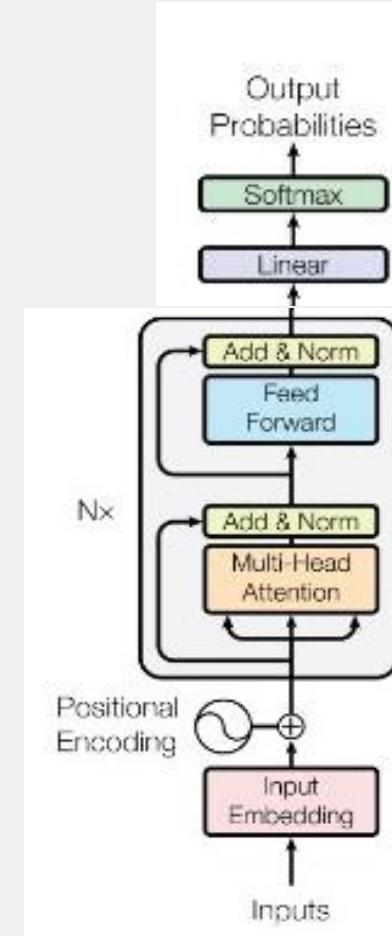
[2] <https://arxiv.org/pdf/2104.04473>



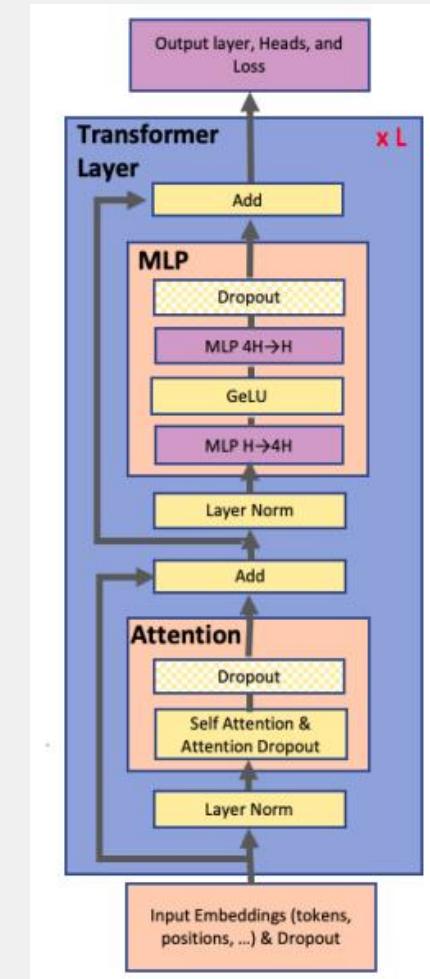
RECAP ON LLM ARCHITECTURE – TRANSFORMER BLOCK



ZOOM IN

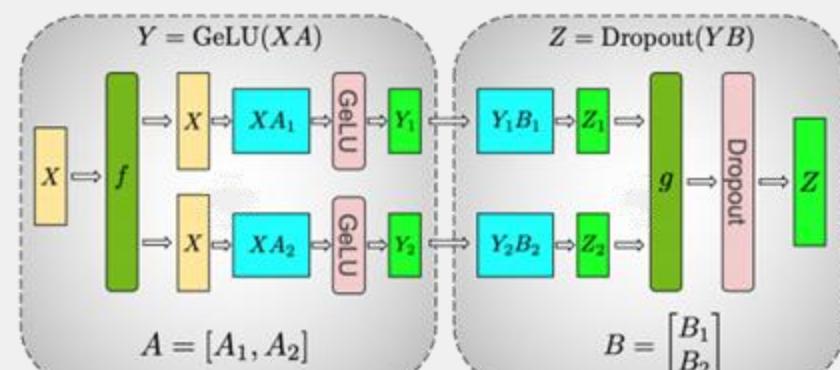
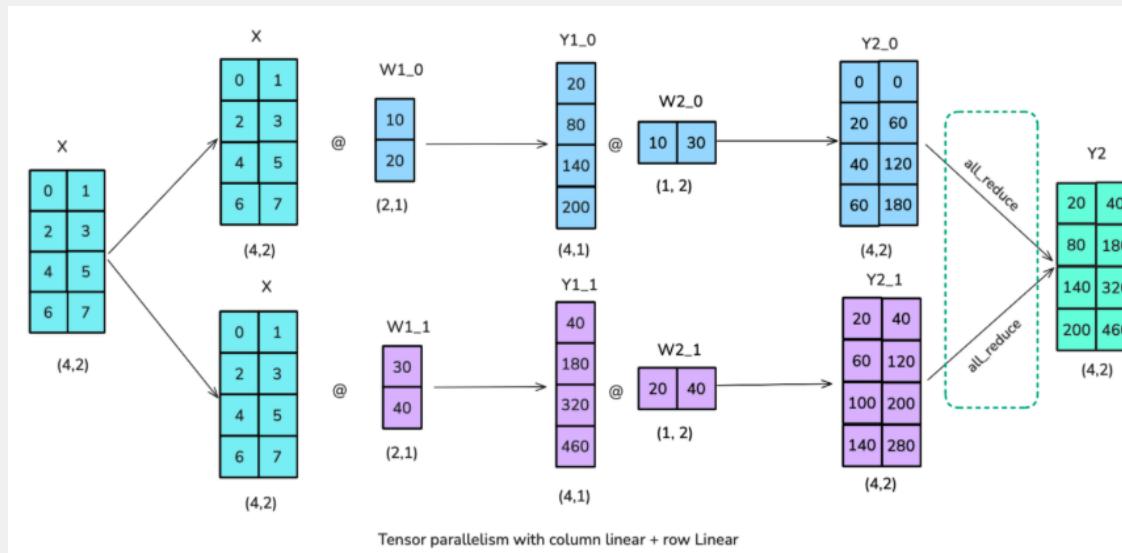


ZOOM IN



TENSOR PARALLELISM IN MLP

The **feedforward** part (2-layer MLP) can be parallelized by having a column-linear followed by a row-linear split. This setup is more efficient than starting with a row-linear followed by column-linear split, as we can skip the intermediate all-reduce between the split operations.

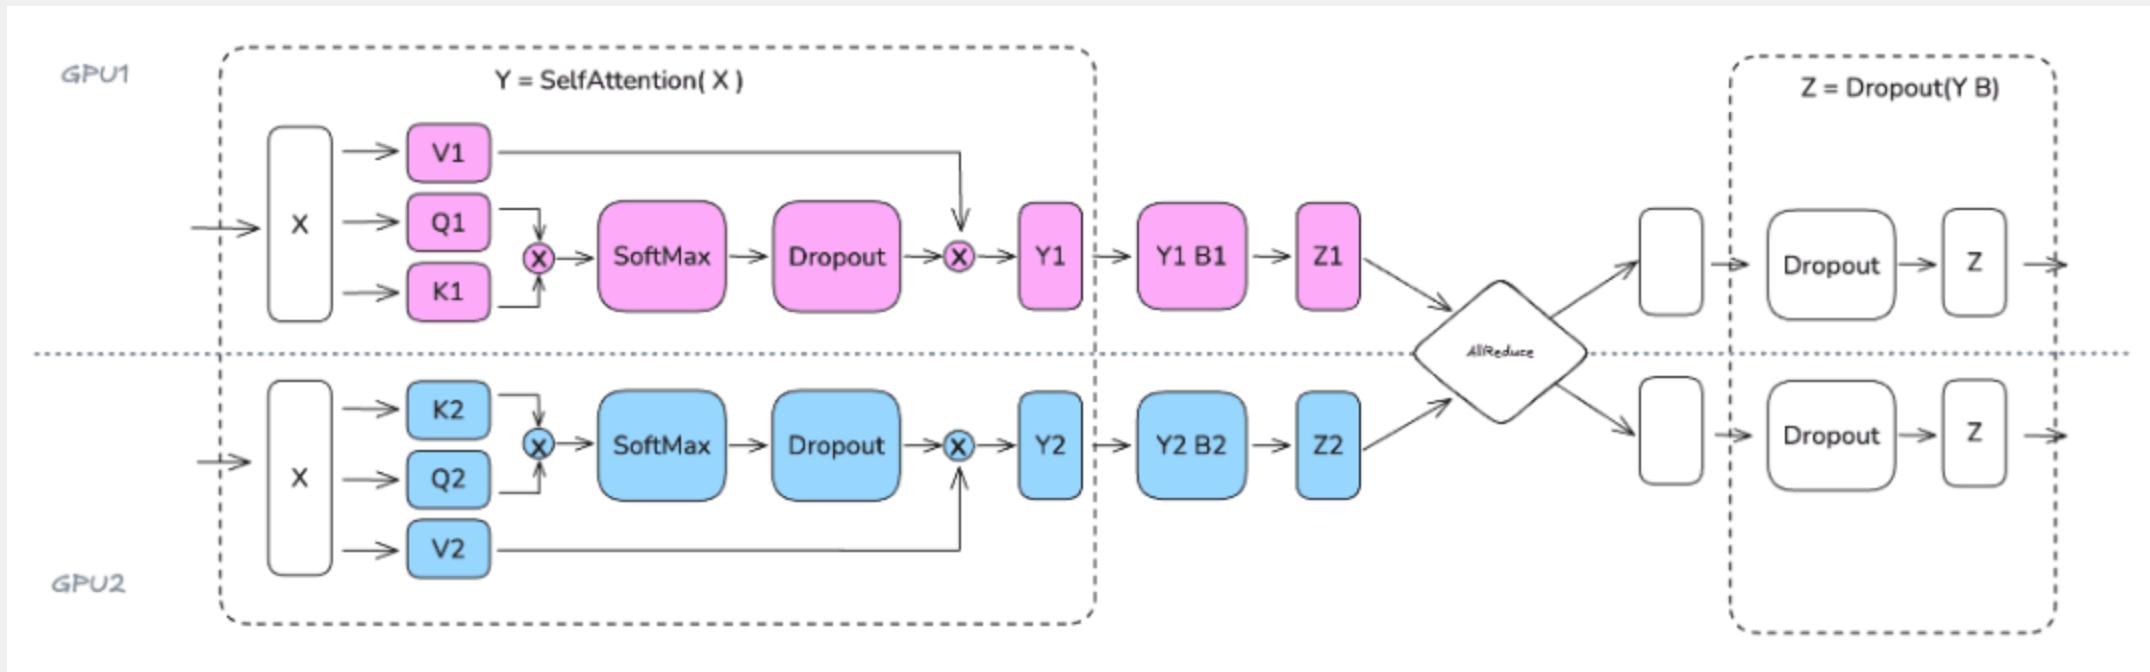


(a) MLP

TENSOR PARALLELISM IN SELF ATTENTION

A similar approach can be used here, where

- Query (Q), Key (K), and Value (V) matrices are split in a **column**-parallel fashion *
- the output projection can be considered a **row**-linear in the GEMM after attention **

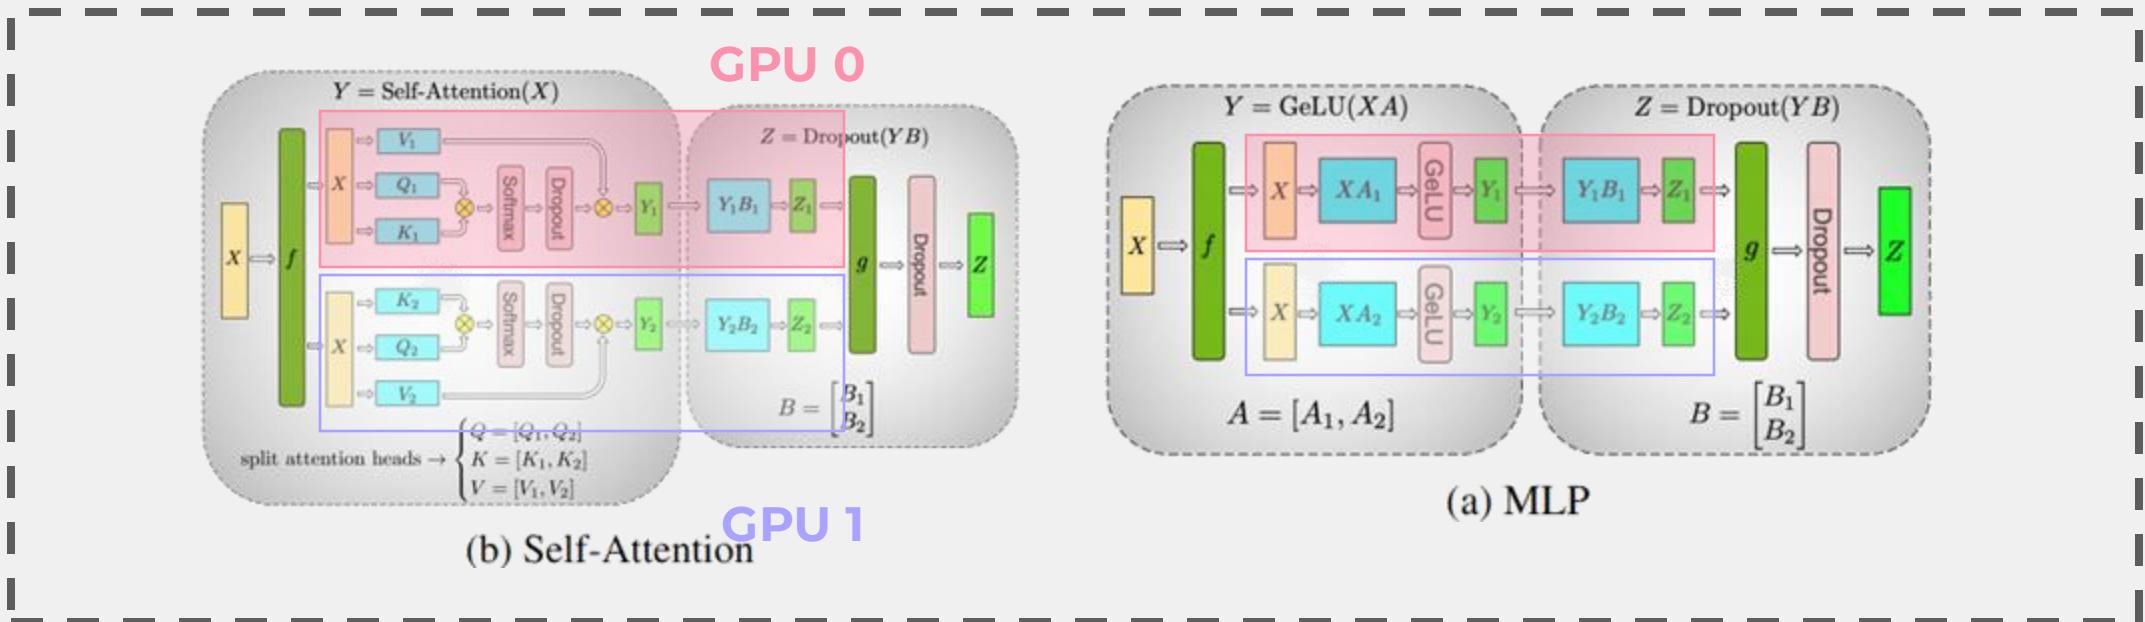


* in multihead attention operation, partitioning K,Q,W in a column parallel fashion results into each attention head is done locally on one GPU.

** this linear layer projects the combined features coming from the attention into a useful representation.

TENSOR PARALLELISM IN A TRANSFORMER BLOCK

Transformer block



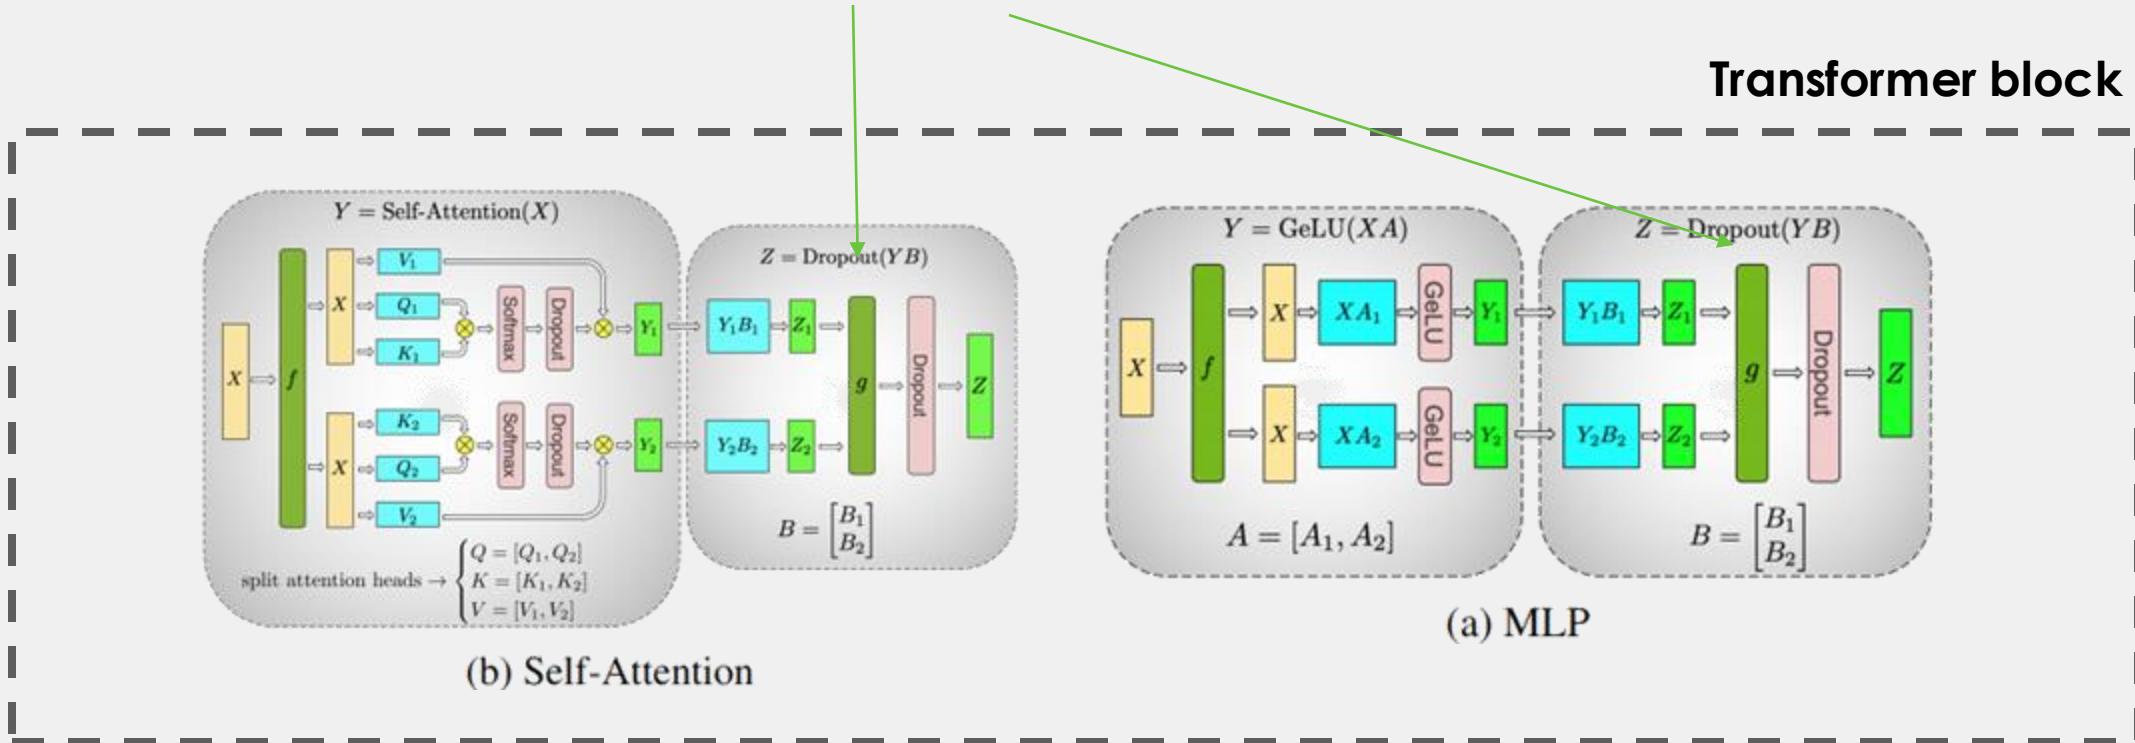
[1] <https://github.com/huggingface/transformers/issues/10321#issuecomment-783543530>

[2] <https://arxiv.org/pdf/2104.04473>

TENSOR PARALLELISM COMM

2 all-reduce in the forward pass

Transformer block

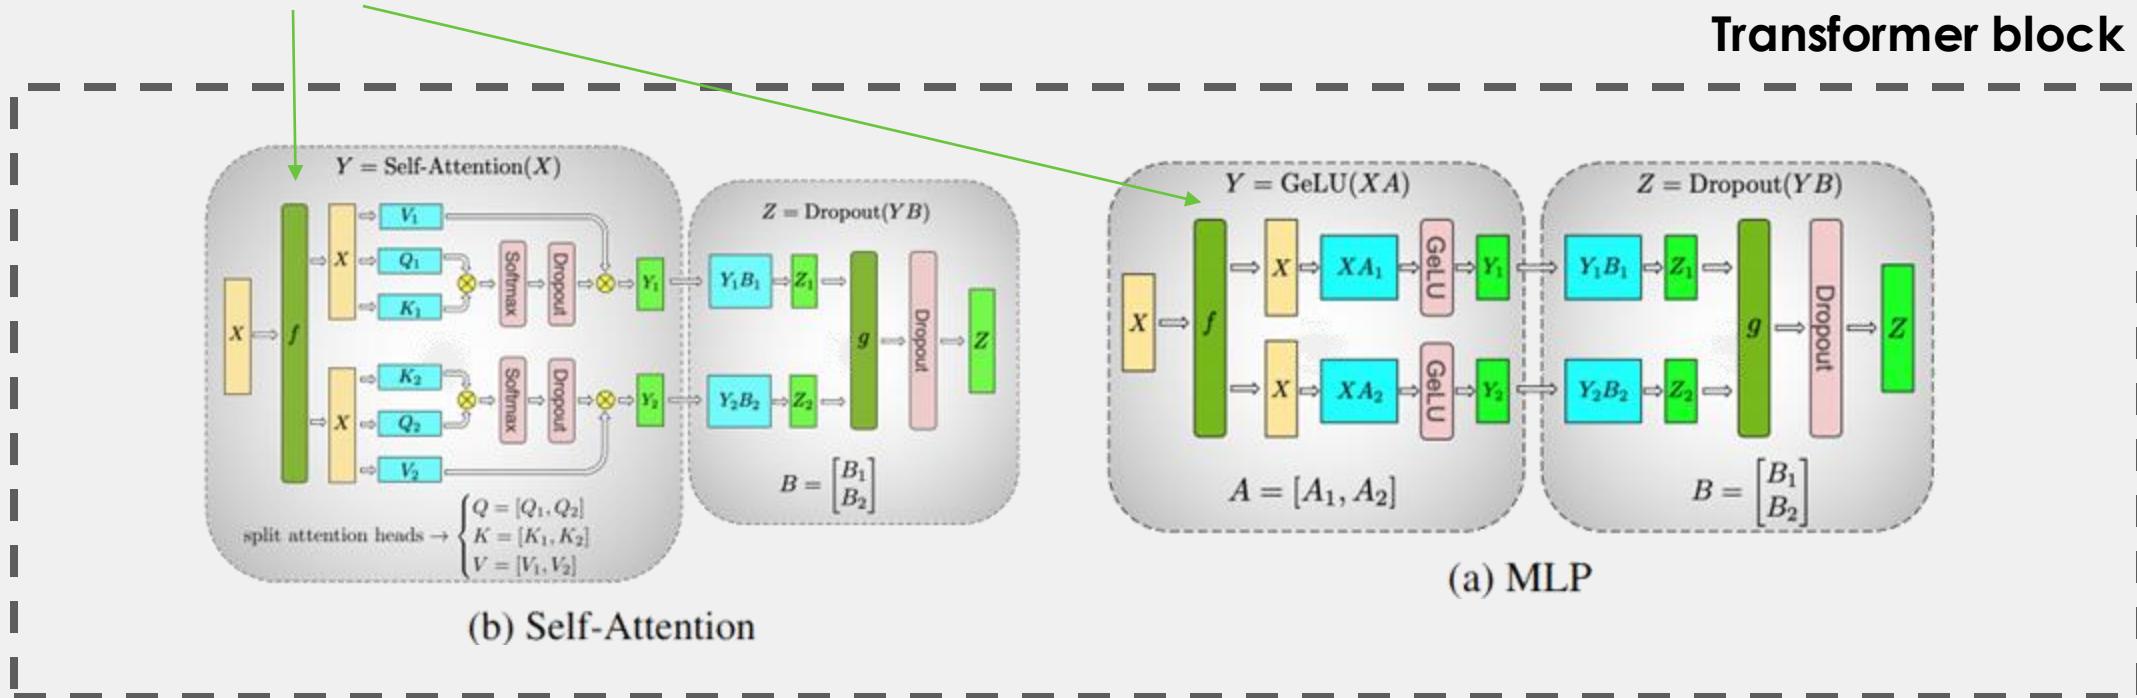


[1] <https://github.com/huggingface/transformers/issues/10321#issuecomment-783543530>

[2] <https://arxiv.org/pdf/2104.04473>

TENSOR PARALLELISM COMM

2 all-reduce in the backward pass

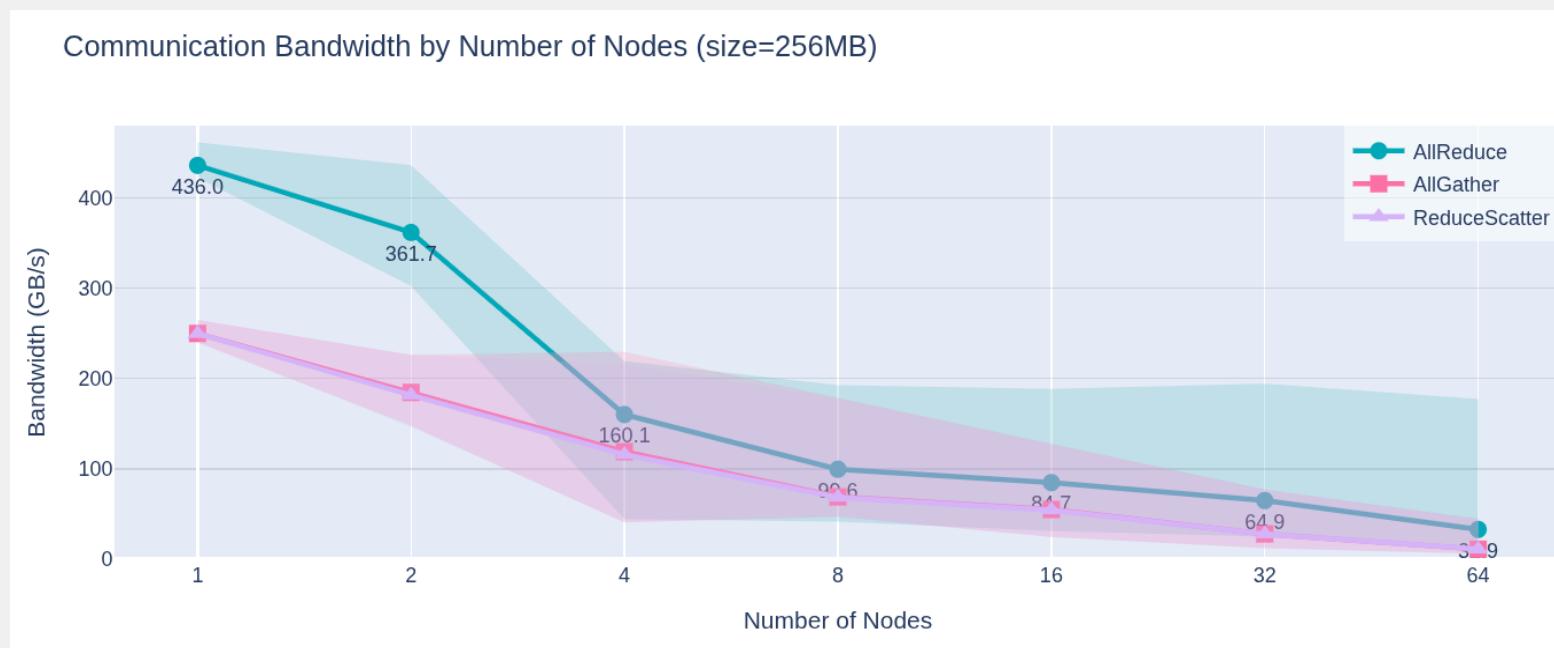


[1] <https://github.com/huggingface/transformers/issues/10321#issuecomment-783543530>

[2] <https://arxiv.org/pdf/2104.04473>

DRAWBACK OF TENSOR PARALLELISM

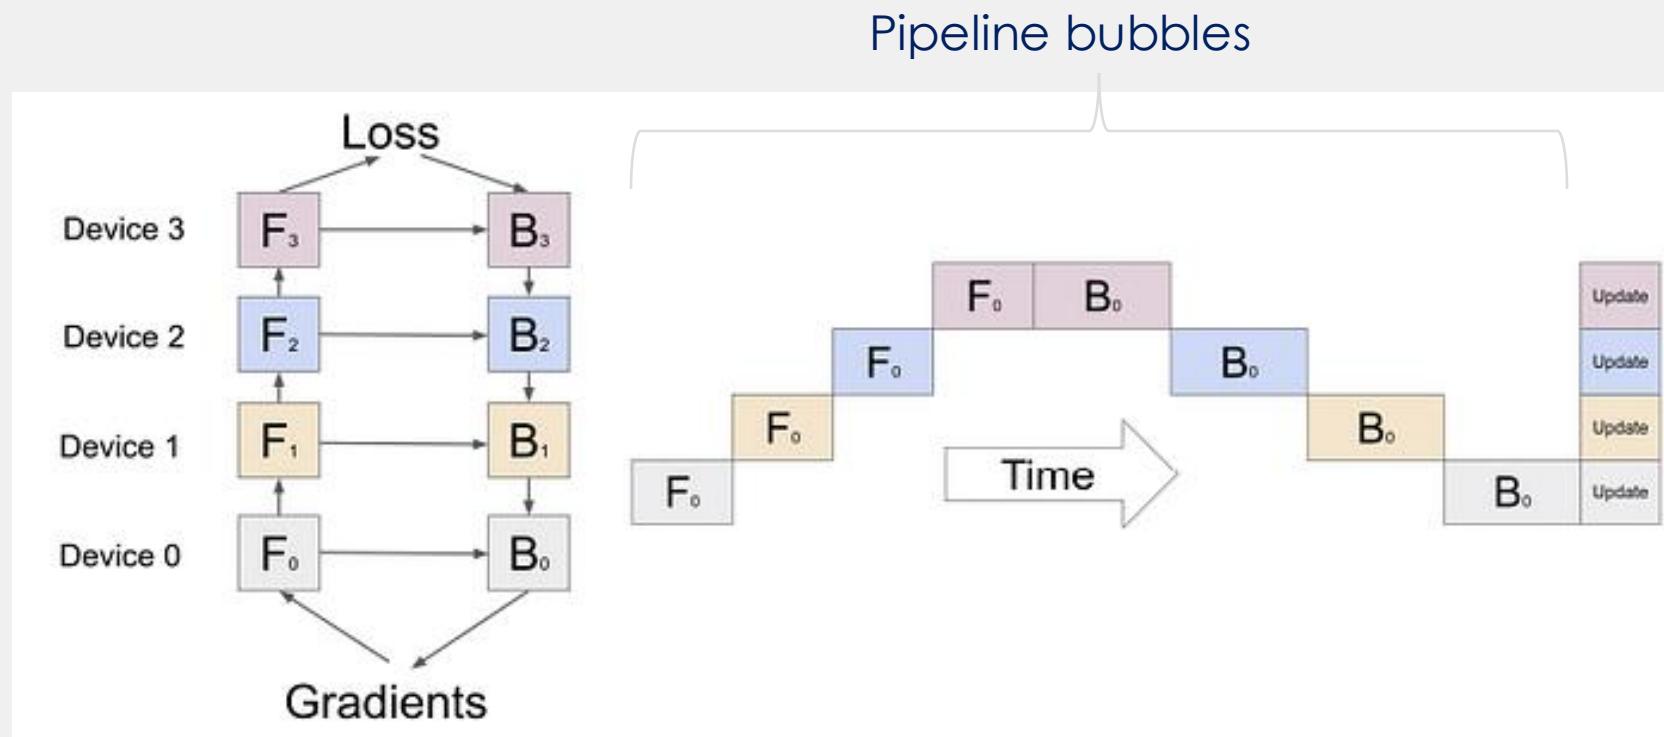
Trying to scale TP past the number of GPUs on a single node ($TP > 4$ in our case) forces to use lower-bandwidth network communication (intra-node), which can significantly impair performance.



We can solve this issue by adding another parallelism dimension: **pipeline parallelism (PP)**.

PIPELINE PARALLELISM

The model is split at layer level on multiple gpus.

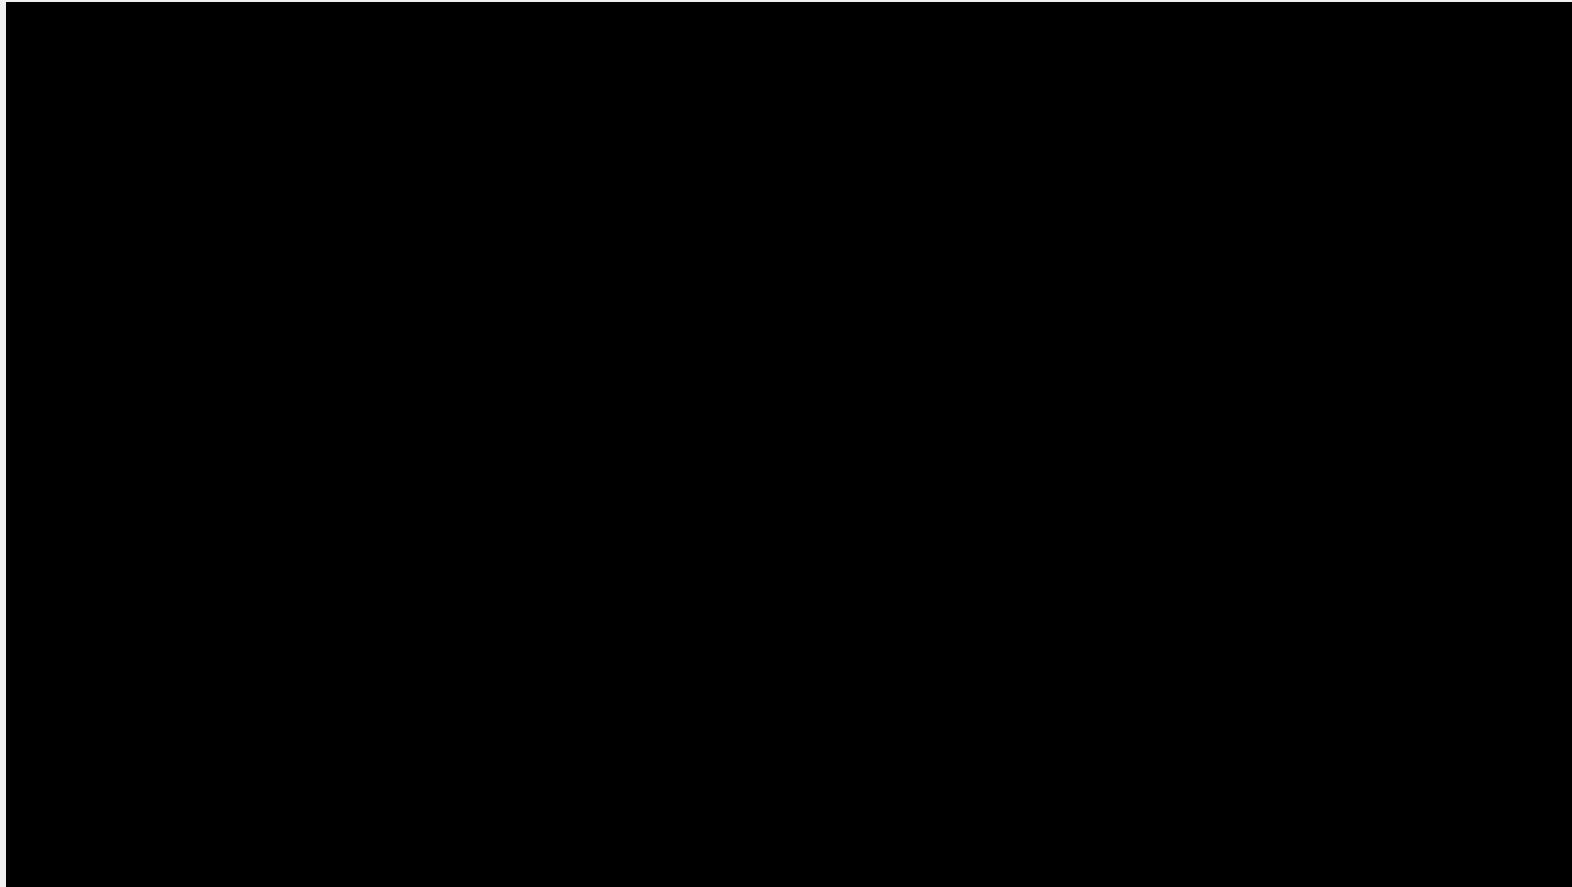


[1] <https://medium.com/byte-sized-ai/pipeline-parallelism-explained-in-2-mins-6bdf1ab29053>

[2] <https://arxiv.org/pdf/2104.04473>



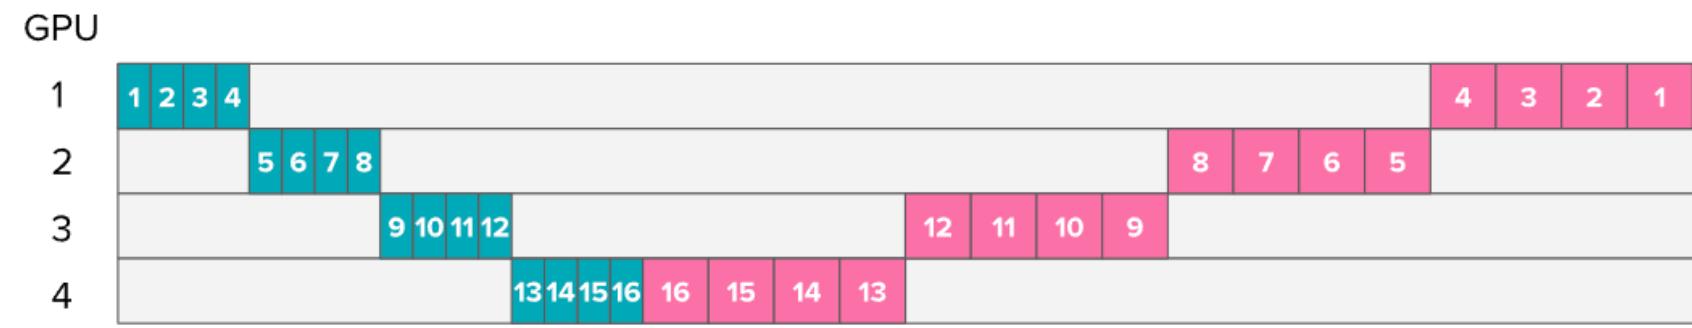
PIPELINE PARALLELISM



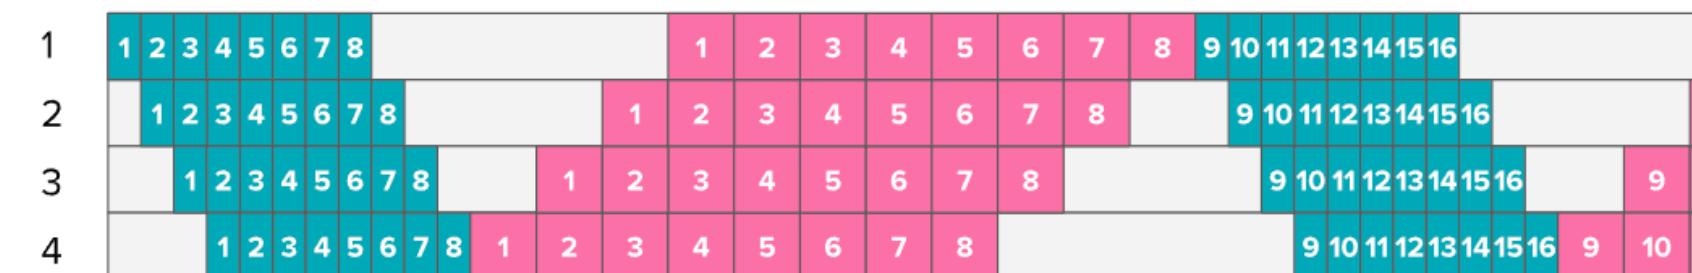
[1] <https://docs.nvidia.com/nemo/megatron-bridge/0.2.0/parallelisms.html>

PIPELINE PARALLELISM

Vanilla



AFAB



1F1B



Time →

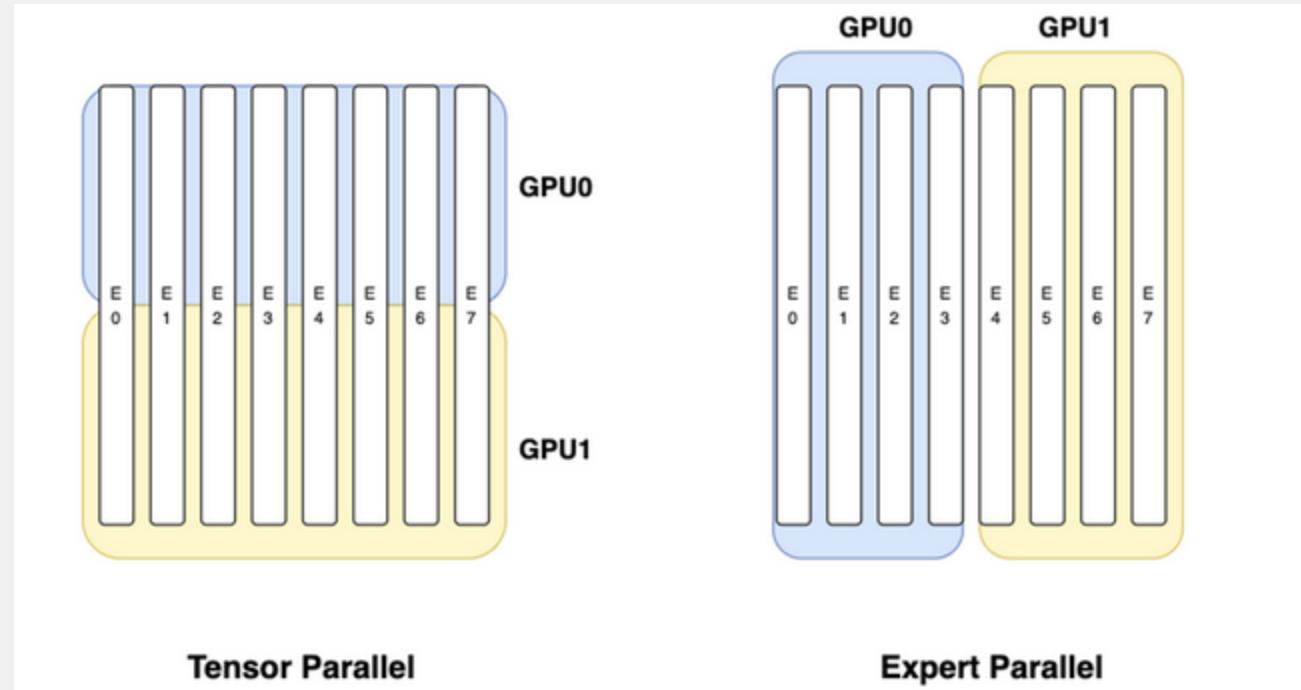
Forward pass

Backward pass

Device idle

fwd: rcv/snd activ.
bwd: rcv/snd grad.

EXPERT PARALLELISM



<https://nvidia.github.io/TensorRT-LLM/advanced/expert-parallelism.html>



HYBRID PARALLELISM

- DP + PP (2D parallelism)

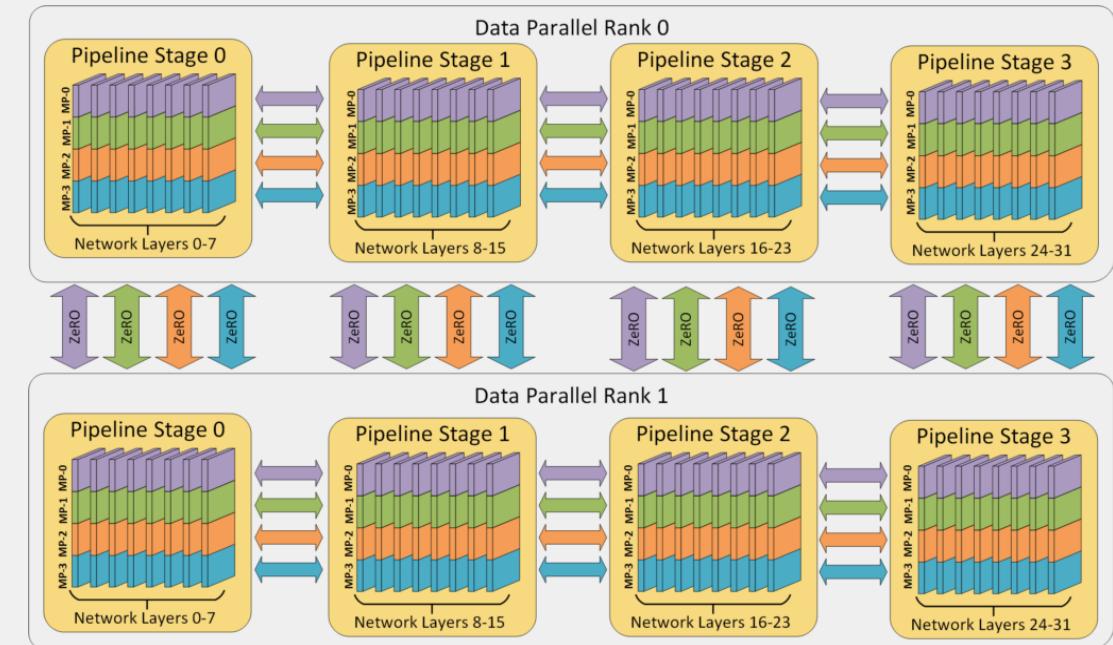
- DP+PP+TP (3D parallelism)

- DP+PP+TP+ZeRO1

- PP + TP + EP

- DP + PP + TP + SP/CP + EP (5D parallelism)

3D parallelism



<https://www.microsoft.com/en-us/research/blog/deepspeed-extreme-scale-model-training-for-everyone/>

Interesting article on 5D parallelism :

https://huggingface.co/spaces/nanotron/ultrascale-playbook?section=high_level_overview



MEGATRON-LM

- Create singularity container
- Convert a model checkpoint in Megatron-LM format (not needed if you train from scratch)
- Tokenize the dataset and create a .pbin file

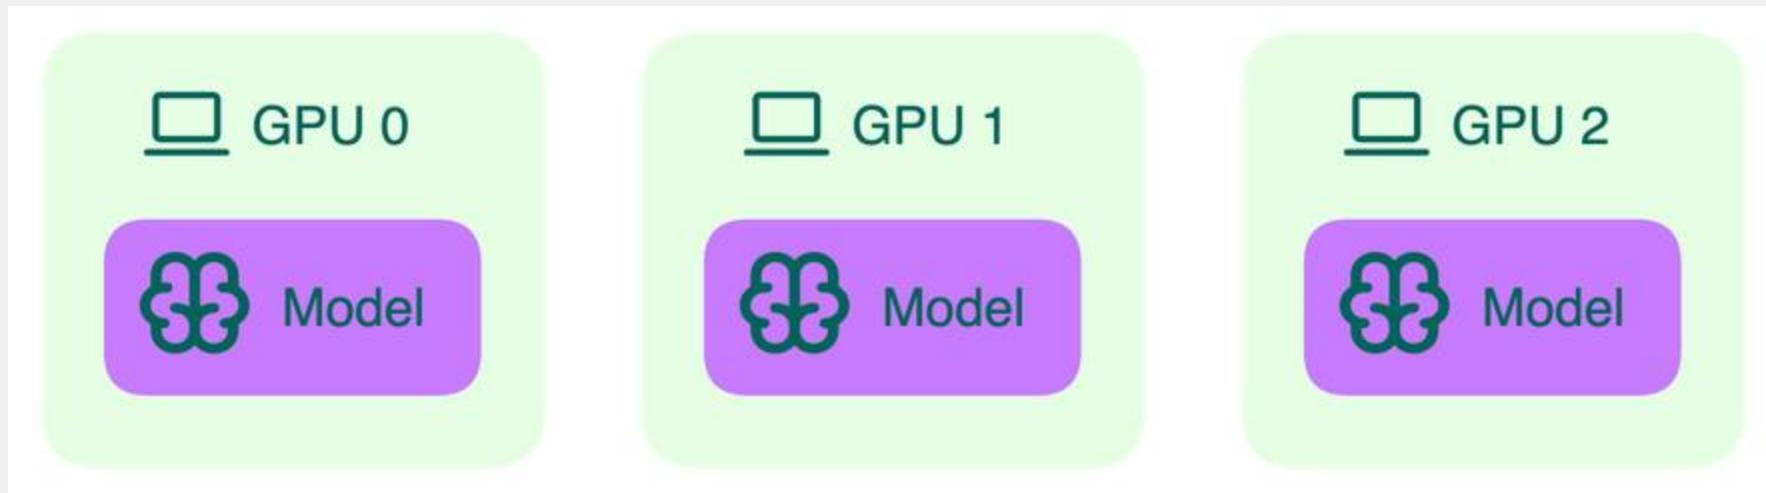


MEGATRON-LM EXERCISE

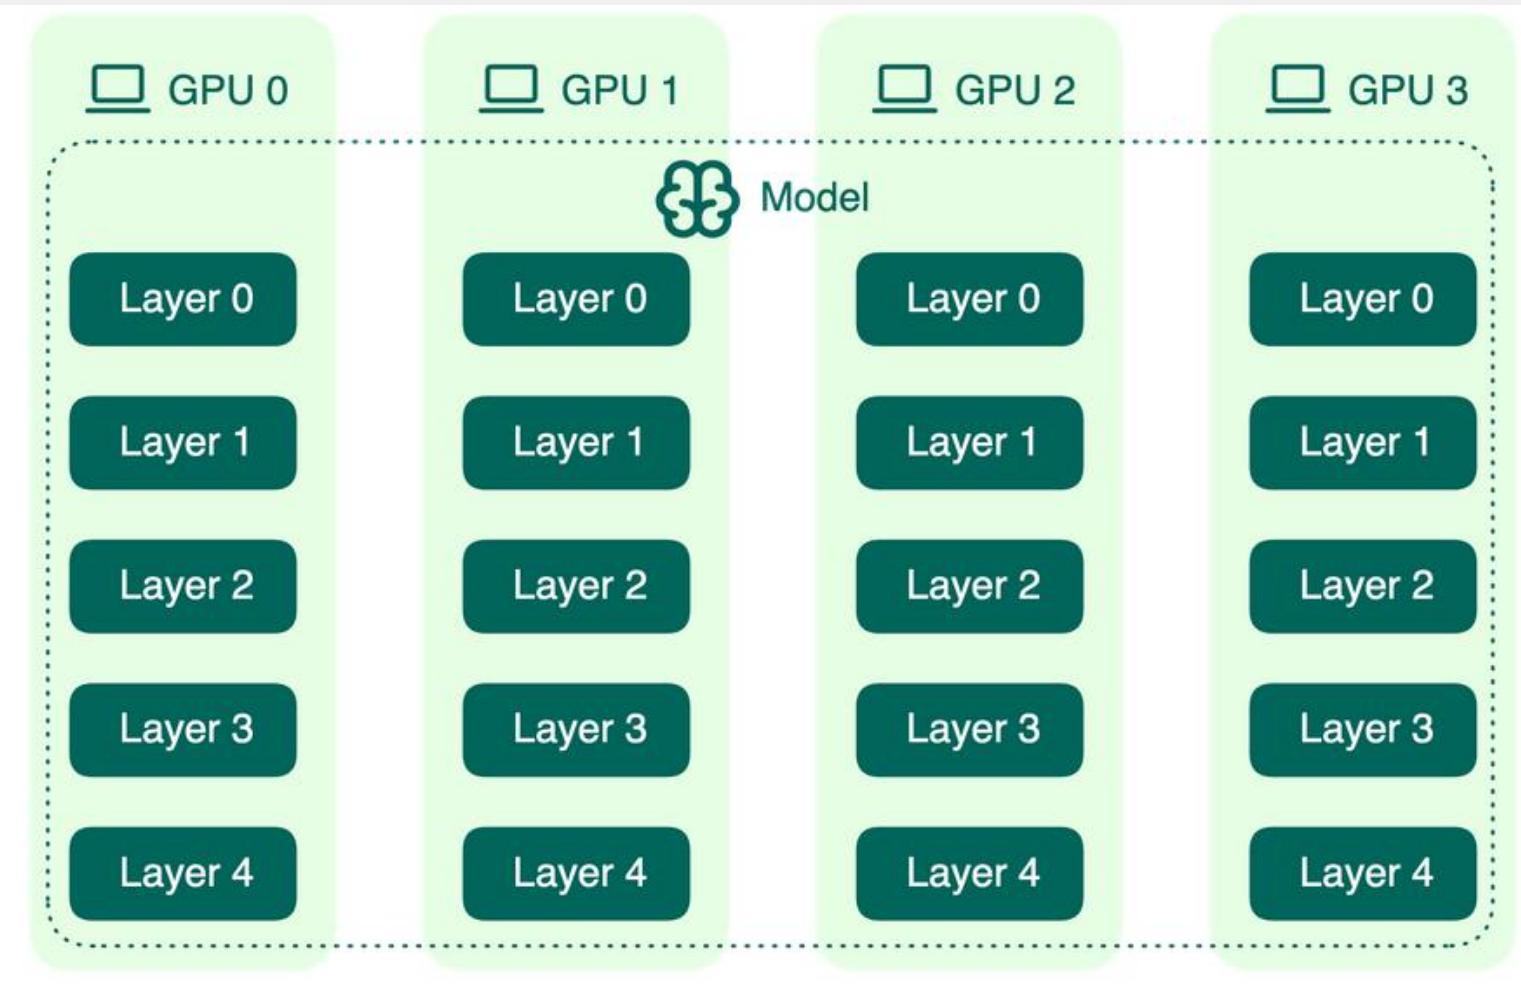
- Create your model architecture (less than 8B params!)
- Complete the jobscript `launch_megatron_LMM.sh`
- Do some tests with different hyperparams
- How did you choose TP and PP?



TEST YOUR KNOWLEDGE

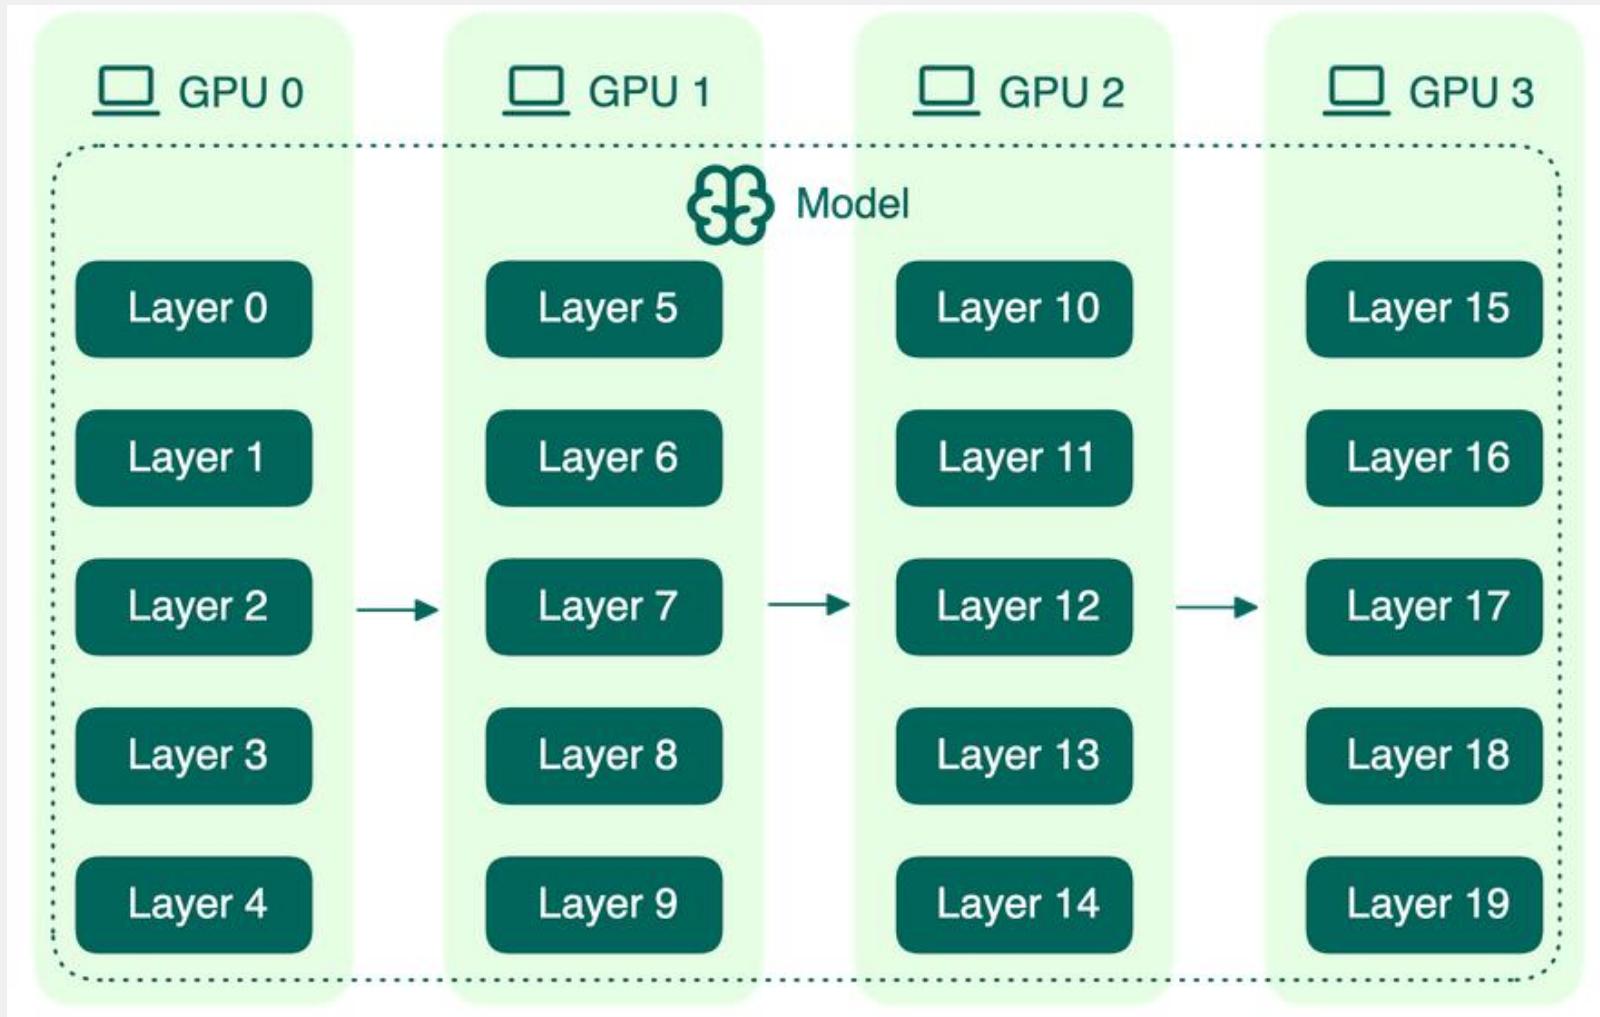


TEST YOUR KNOWLEDGE

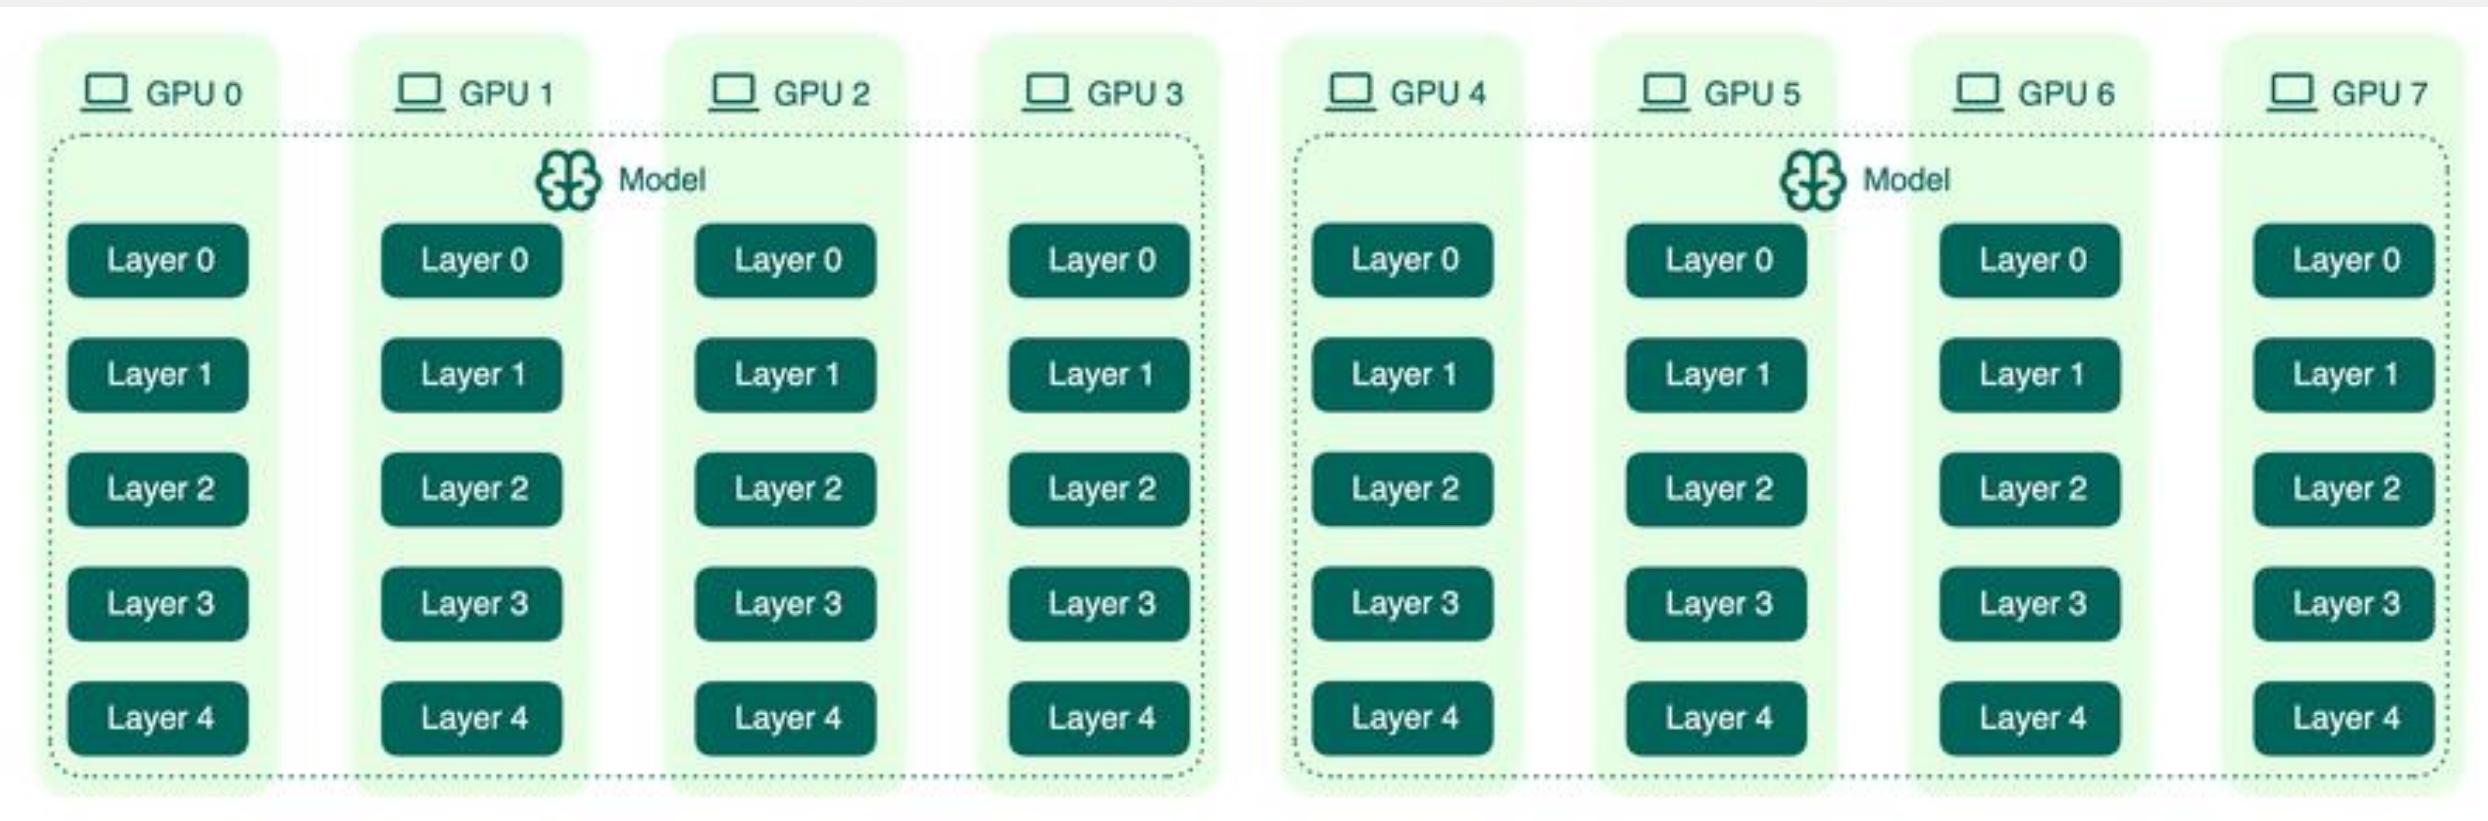




TEST YOUR KNOWLEDGE



TEST YOUR KNOWLEDGE





RESULTS – OLMO-CORE

Size	#gpus	#replicas	GBS	MBS*	parallelism	Avg Tflops	Max Tflops
1B	8	2	256	4	HSDP	174	186
1B	16	2	512	4	HSDP	142	190
1B	32	2	1024	4	HSDP	140	190
8B	8	1	512	2	FSDP	118	119
8B	16	1	1024	2	FSDP	117	118
8B	32	1	2048	2	FSDP	112	113
8B	8	2	512	1	HSDP	-----	-----
8B	16	2	1024	1	HSDP	82	81
8B	32	2	2048	1	HSDP	80	80
8B	32	2	512	2	HSDP	101	104

*This must evenly divide into the global batch size by a factor of the data parallel world size. If this is less than the global batch divided by the data parallel world size then gradient accumulation is used.



SCALING HSDP

Size	#gpus	#replicas	GBS	MBS	parallelism	Avg Tflops	Max Tflops
1B	8	2	256	4	HSDP	174	186
1B	16	2	512	4	HSDP	142	190
1B	32	2	1024	4	HSDP	140	190
8B	8	1	512	2	FSDP	118	119
8B	16	1	1024	2	FSDP	117	118
8B	32	1	2048	2	FSDP	112	113
8B	8	2	512	1	HSDP	-----	----
8B	16	2	1024	1	HSDP	82	81
8B	32	2	2048	1	HSDP	80	80
8B	32	2	512	2	HSDP	101	104



HSDP MEMORY ISSUES

Size	#gpus	#replicas	GBS	MBS	parallelism	Avg Tflops	Max Tflops
1B	8	2	256	4	HSDP	174	186
1B	16	2	512	4	HSDP	142	190
1B	32	2	1024	4	HSDP	140	190
8B	8	1	512	2	FSDP	118	119
8B	16	1	1024	2	FSDP	117	118
8B	32	1	2048	2	FSDP	112	113
8B	8	2	512	1	HSDP	-----	----
8B	16	2	1024	1	HSDP	82	81
8B	32	2	2048	1	HSDP	80	80
8B	32	2	512	2	HSDP	101	104



MEGATRON-LM : How to keep constant Tflops

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	1	144	147
8B	32	4	2	512	2	149	152
8B	64	4	2	512	2	121	136
8B	128	4	2	512	2	109	113
8B	256	4	2	512	1	79	83
8B	256	4	2	1024	1	100	109



MEGATRON-LM : How to keep constant Tflops

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	1	144	147
8B	32	4	2	512	2	149	152
8B	64	4	2	512	2	121	136
8B	128	4	2	512	2	109	113
8B	256	4	2	512	1	79	83
8B	256	4	2	1024	1	100	109



MEGATRON-LM : Microbatch size matters

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	1	144	147
8B	32	4	2	512	2	149	152
8B	64	4	2	512	2	121	136
8B	128	4	2	512	2	109	113
8B	256	4	2	512	1	79	83
8B	256	4	2	1024	1	100	109



MEGATRON-LM : Microbatch size matters

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	1	144	147
8B	32	4	2	512	2	149	152
8B	64	4	2	512	2	121	136
8B	128	4	2	512	2	109	113
8B	256	4	2	512	1	79	83
8B	256	4	2	1024	1	100	109



MEGATRON-LM : Larger model, higher PP

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	2	149	152
32B	64	4	16	512	1	110	115
32B	64	4	16	1024	1	114	116
32B	128	4	16	1024	1	110	111
70B	128	4	16	1024	1	96	114



MEGATRON-LM : Larger models, lower tflops

Size	#gpus	TP	PP	GBS	MBS	Avg Tflops	Max Tflops
8B	8	4	2	512	1	160	162
8B	16	4	2	512	1	155	158
8B	16	4	2	512	2	160	163
8B	32	4	2	512	2	149	152
32B	64	4	16	512	1	110	115
32B	64	4	16	1024	1	114	116
32B	128	4	16	1024	1	110	111
70B	128	4	16	1024	1	96	114

Reach out to us @

info@minerva4ai.eu

Thank you



**Co-funded by
the European Union**



This project has received funding from the European High-Performance Computing Joint Undertaking (JU) under grant agreement No 101182737. The JU receives support from the Digital Europe Programme.